


# Operating Systems

---

**Nguyen Tri Thanh**  
**ntthanh@vnu.edu.vn**



# Question



What is **incorrect** about overlays?

- A. overlays allows a large program to run in a smaller MEM
- B. Overlays only loads codes on demand (when they are used)
- C. Programmers need to split the program into modules
- D. Overlays is supported in all high level programming languages

# Question



What is **incorrect** about swapping?

- A. swapping is the same as overlays
- B. swapping uses hard disk as the *backing store*
- C. swapping allows many processes whose size is even larger than MEM to run
- D. a lower priority process is rolled out for a higher priority one to run (when needed)

# Review



Which is incorrect about non-contiguous MEM allocation?

- A. split logical memory into parts
- B. utilize MEM more effectively in comparison with contiguous allocation method
- C. need a Memory Management Unit
- D. only suitable for some types of processes

# Review



Which is correct about MMU of paging and segmentation allocation methods?

- A. they are the same
- B. MMU of paging needs more information than that of segmentation
- C. they use different resolution methods
- D. MMU of segmentation is faster than that of paging



## Question

Suppose a process in contiguous allocation:

- the base address is 10400
- the limit register is 1200
- the reference is 246;

Which of the following is the correct physical address of the reference ?

- A. 10154
- B. 10646
- C. 1446
- D. 954

# Question




A system uses paging

- the frame size of 2KB;
- the address register is 32 bits

Which of the following is correct about register segmentation?

- A. (page:offset) = (19:13)
- B. (page:offset) = (21:11)
- C. (page:offset) = (22:10)
- D. (page:offset) = (20:12)

# Question



Frame
56
120
3

A system uses paging

- the frame size of 4KB;
- the address register is 32 bits
- Which of the following is the correct physical address of the reference (2,1296)?
  - A.  $560 \times 4096 + 1296$
  - B.  $120 \times 4096 + 1296$
  - C.  $3 \times 4096 + 1296$
  - D.  $120 \times 1024 + 1296$



# Virtual Memory


---

Paging on demand

Page replacement

Frame allocation

Thrashing





# Objectives

- Introduce paging method
- Introduce segmentation method

# Reference



- Chapter 9 of **Operating System Concepts**



# Virtual memory

# Virtual memory



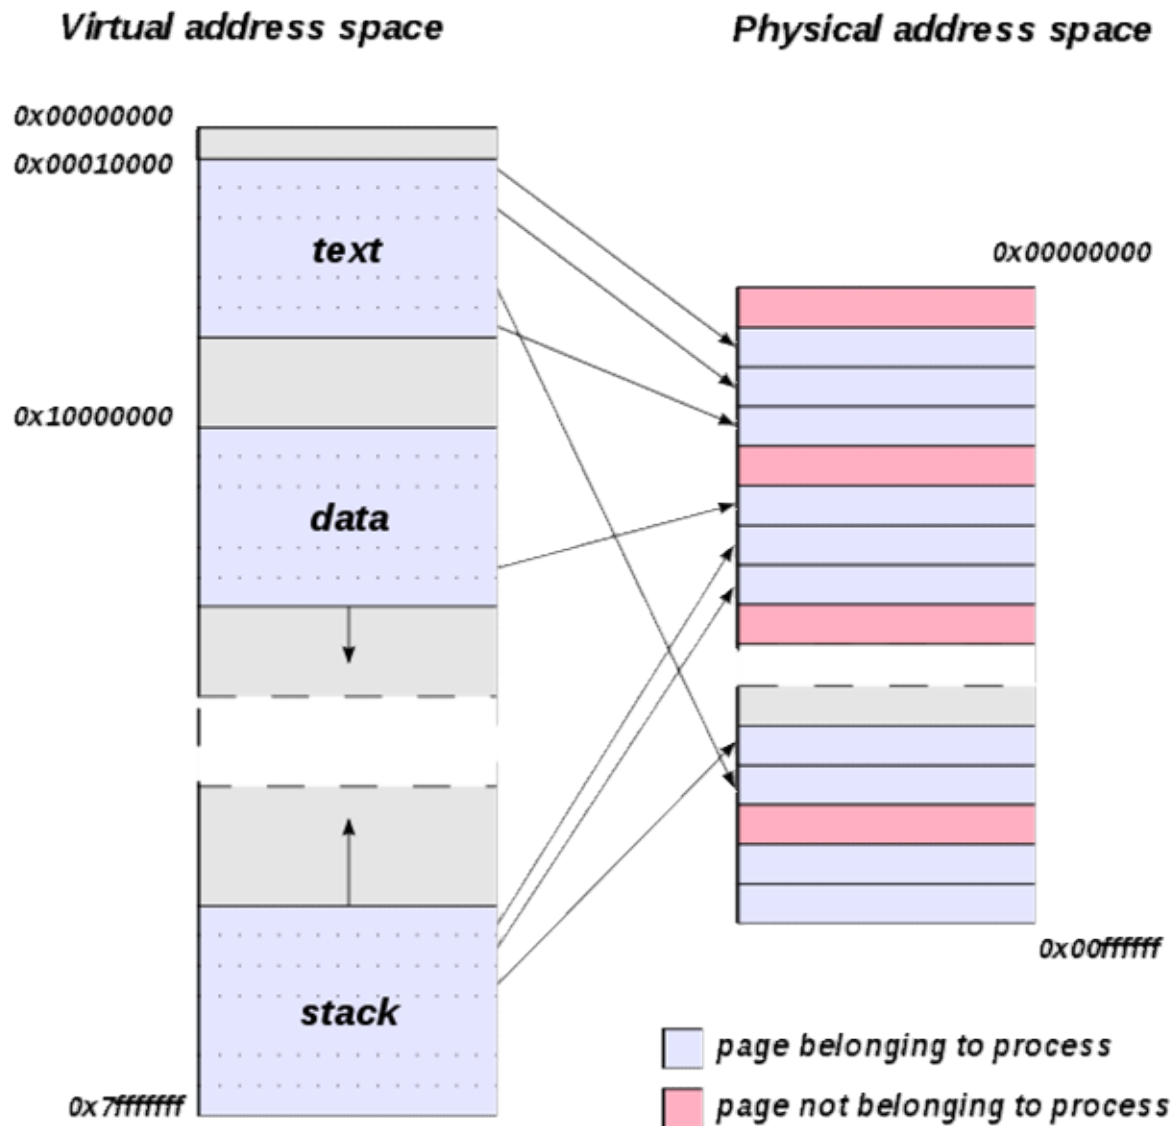
- Separation of user logical memory from physical memory.
  - Only a **part** of the program needs to be in memory for execution
  - Logical address space can therefore be much **larger** than physical address space
  - Allows address spaces to be **shared** by several processes
  - Allows for more efficient process creation
- Virtual memory can be **implemented** via
  - Paging on demand
  - Segmentation on demand



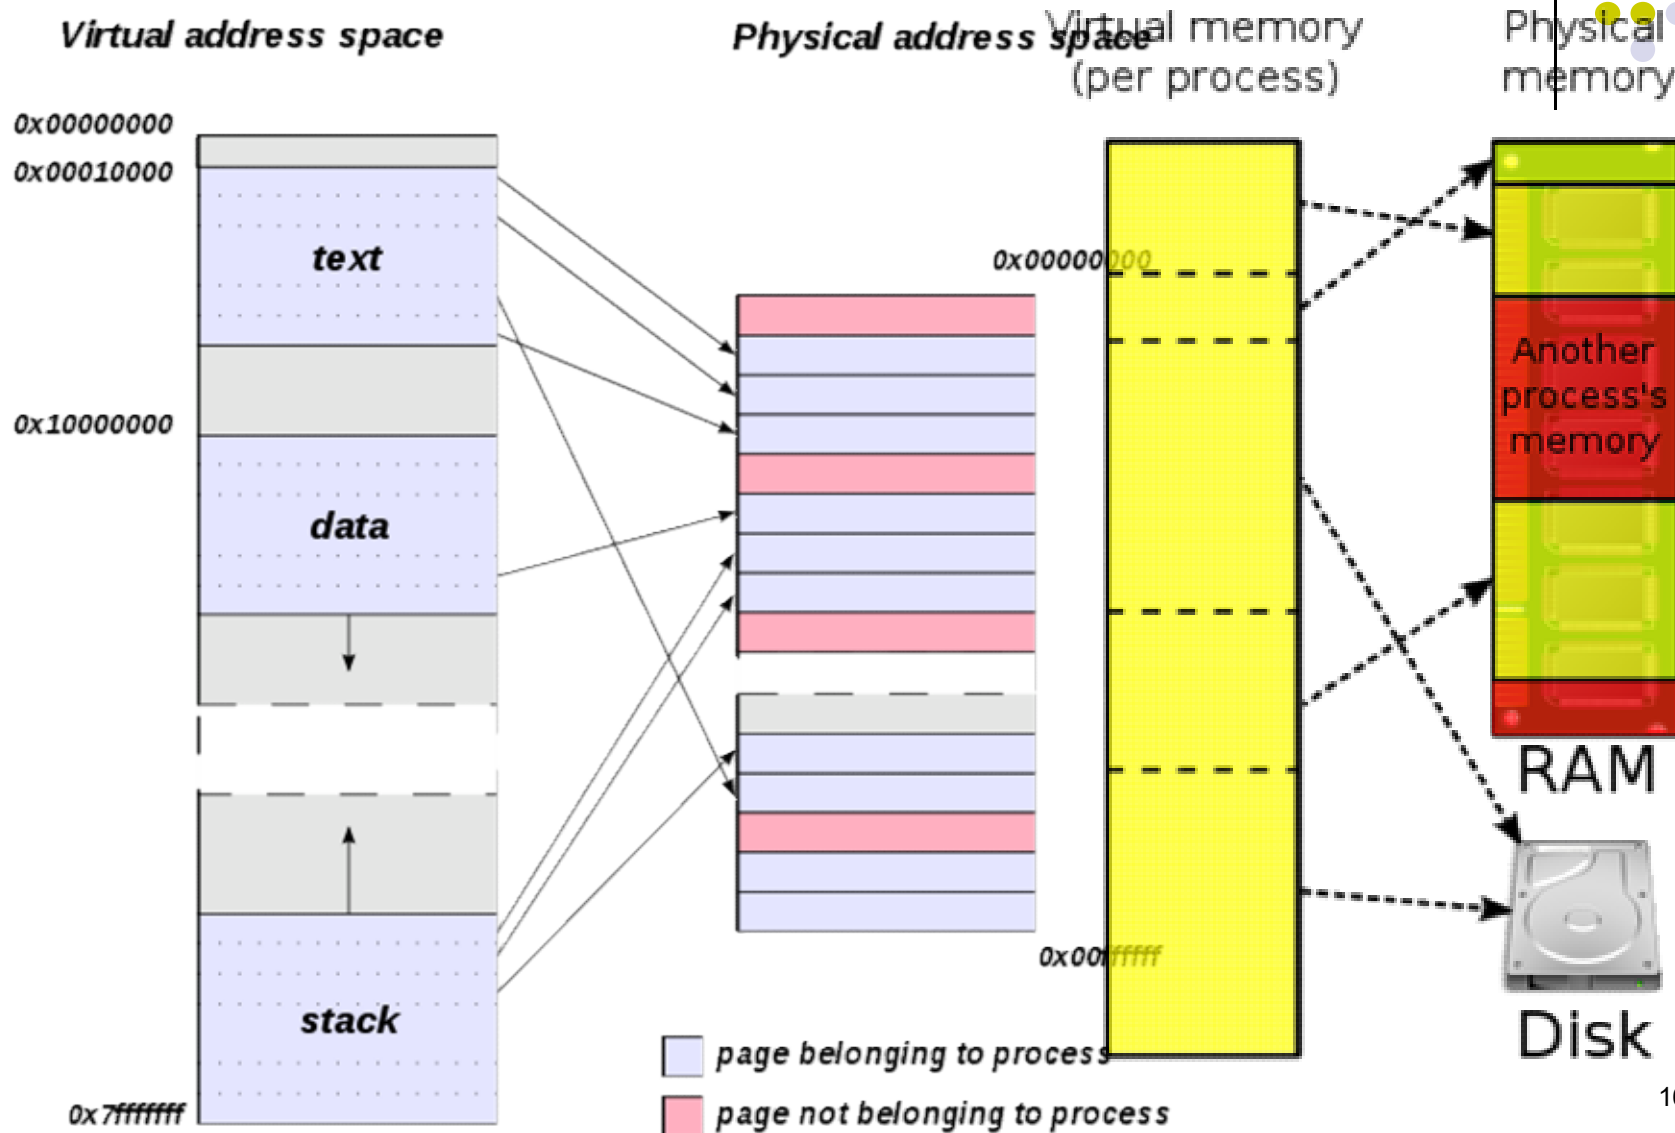
# Virtual memory

- Linux is, of course, a **virtual memory system**, meaning that the addresses seen by user programs do not directly correspond to the physical addresses used by the hardware. Virtual memory introduces a layer of indirection that allows a number of nice things. With virtual memory, programs running on the system can allocate **far more memory than is physically available**; indeed, even a single process can have a virtual address space larger than the system's physical memory. Virtual memory also allows the program to play a number of tricks with the process's address space, including mapping the program's memory to device memory.
- <http://www.makelinux.net/ldd3/chp-15-sect-1>

# Virtual-address Space




# Virtual-address Space

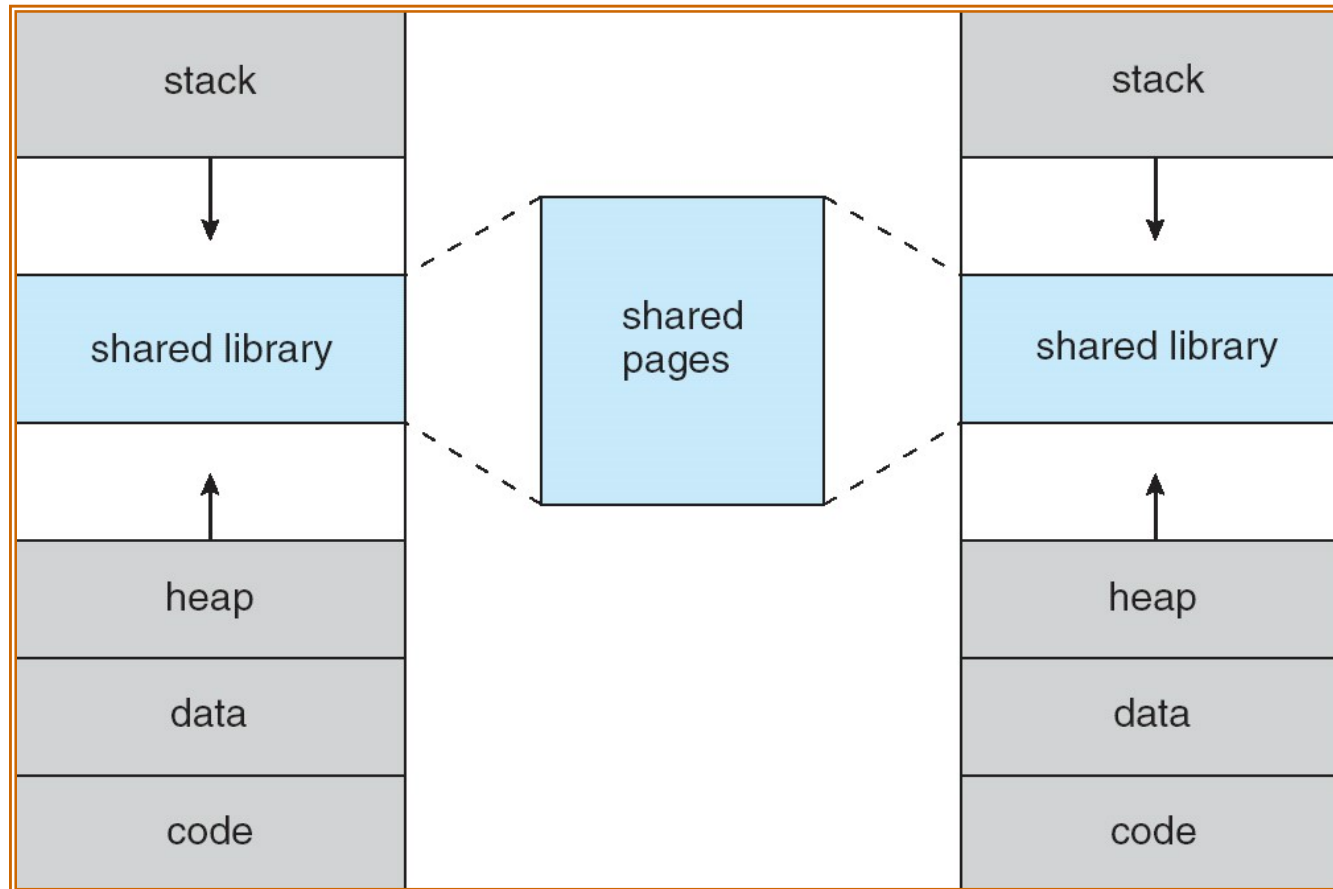
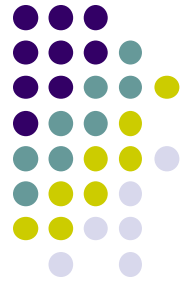




# Virtual Memory That is Larger Than Physical Memory



# Shared Library Using Virtual Memory





# Process Creation

- Virtual memory allows other benefits
  - Copy-on-Write during process creation
  - Memory-Mapped Files

# Copy-on-Write



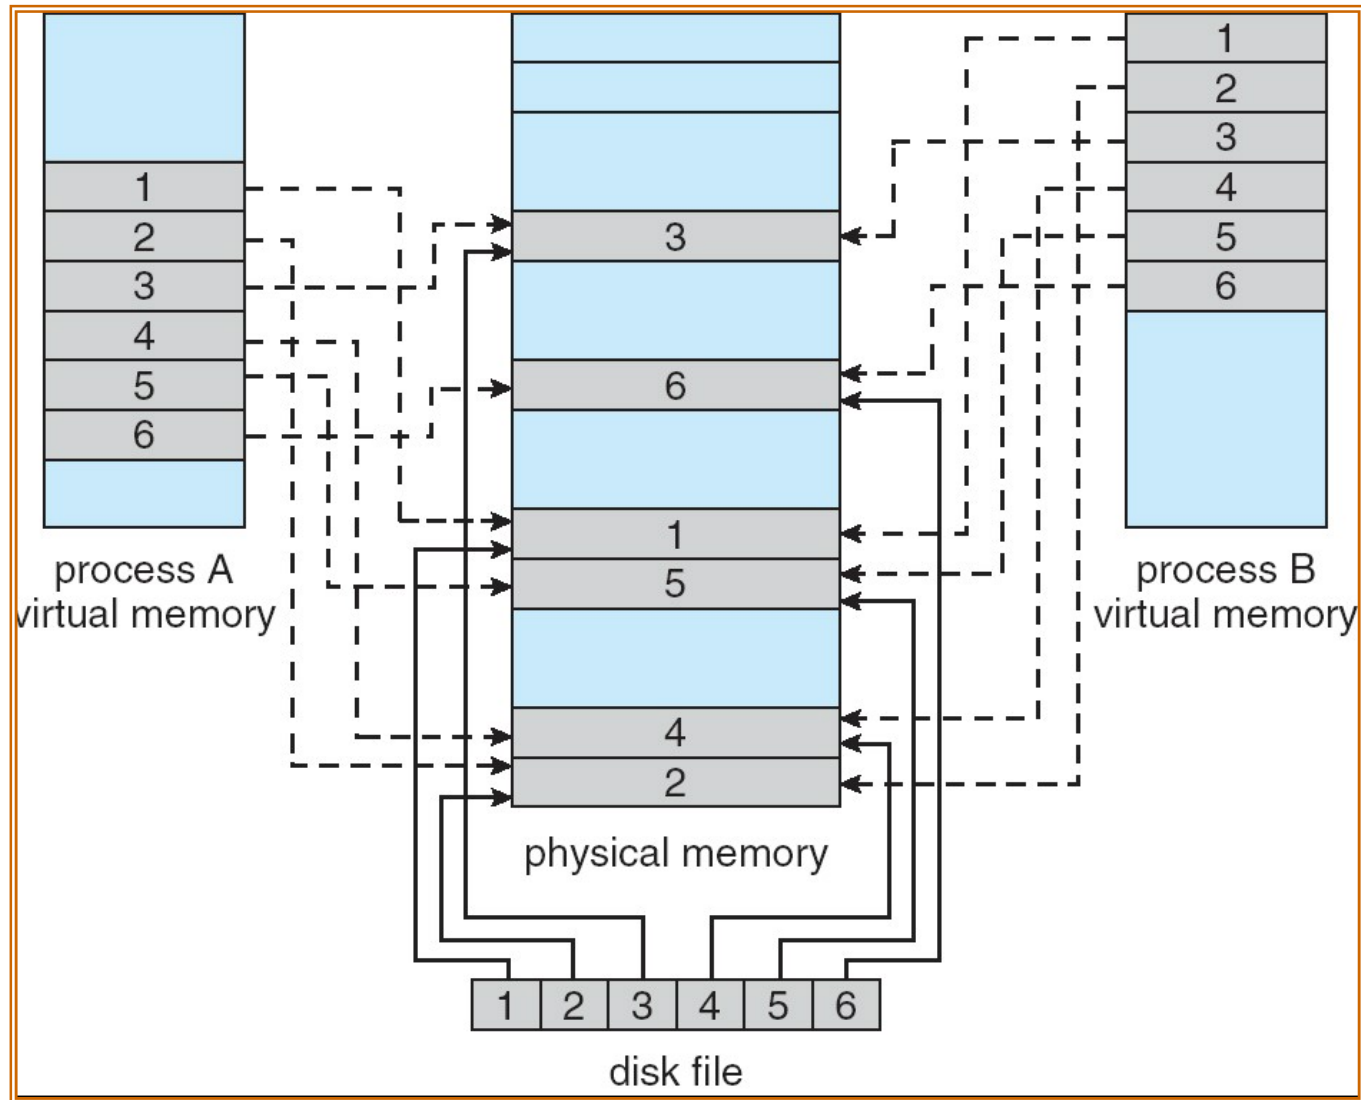
- Copy-on-Write (COW) allows both parent and child processes to initially *share* the same pages in memory
  - If either process modifies a shared page, only then is the page copied
  - refer to Sect. 2, Chapter 3 of “Lập trình C/C++ ...”
- COW allows more efficient process creation as only modified pages are copied

# Memory Mapped Files



- A file is considered as a memory segment
- Read/write operations are performed via memory
  - not read/write file system calls
- Allow multiple processes shared a file
- refer to Sect. 5, Chapter 4 of “Lập trình /C++ trên Linux”

# Memory Mapped Files

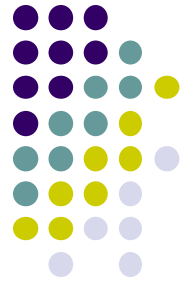




## Question

- What is the correct advantage of memory mapped file?
  - A. reduces the task of the system's OS
  - B. treats as the buffer for manipulating the file
  - C. allows programmers to organize the file
  - D. uses as shared resource among processes

# Question



- Which of the following is incorrect about virtual memory?
  - A. it is separated from physical memory
  - B. it is mapped into physical memory during process execution
  - C. it gives additional benefits, e.g., COW, file mapping
  - D. an address in virtual memory is preserved when mapped into physical memory



# Dynamic loading



- Routine is not loaded until it is called
- Better memory-space utilization
  - unused routine is never loaded
- Useful when
  - large amounts of code are needed to handle infrequently occurring cases
- No special support from the operating system
  - refer to Section 3.4, Chapter 7 of “Lập trình C/C++ ...”

# Dynamic linking and shared library




- Linking postponed until execution time
- Small piece of code, *stub*, used to locate the appropriate memory-resident library routine
  - Stub replaces itself with the address of the routine, and executes the routine
  - Operating system needed to check if routine is in processes' memory address
- Dynamic linking is particularly useful for libraries
- Also known as **shared libraries** in Linux
  - refer to Sect. 3, Chapter 7 of “Lập trình C/C++ trên Linux”




# Paging on demand

# Food order



# Food order




[art\\_links@yahoo.com.vn](mailto:art_links@yahoo.com.vn)






# Food order





# Food order



D1145-001.JPG



D1145-002.JPG



D1145-003.JPG



D1145-004.JPG



D1145-005.JPG



D1145-006.JPG



D1145-007.JPG



D1145-008.JPG



D1145-009.JPG



D1145-010.JPG



D1145-011.JPG



D1145-021.JPG



D1145-031.JPG



D1145-041.JPG



D1145-051.JPG



D1145-061.JPG



D1145-071.JPG



D1145-020.JPG



D1145-030.JPG



D1145-040.JPG



D1145-050.JPG



D1145-060.JPG



o.com.vn



Each person prefers different food




art\_links@yahoo.com.vn





# Food order



Each person prefers different food




One more steak

art\_links@yahoo

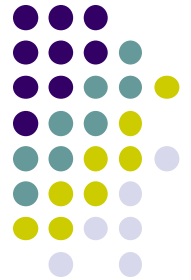


# Food order



Each person prefers different food

# Food order

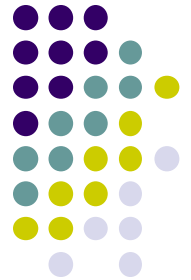


Each person prefers different food





# Food order




One more steak



Each person prefers different food



# Paging on demand




- Bring a page into memory only when needed
  - Less I/O needed
  - Less memory needed
  - Faster response
  - More users
- Page is needed  $\Rightarrow$  reference to it
  - invalid reference  $\Rightarrow$  abort
  - not-in-memory  $\Rightarrow$  bring to memory
- **Lazy swapper** – never swaps a page into memory unless page is needed
  - Swapper that deals with pages is a **pager**

# Question



- Why it is possible when only a part of a process is loaded into MEM?
  - A. Because instructions of a process are independent
  - B. Because we can indicate which instructions to run
  - C. Because only one instruction is executed at a time
  - D. Because related instructions are always in the same group

# Transfer of a Paged Memory to Contiguous Disk Space





# Valid-Invalid Bit

- With each page table entry a valid–invalid bit is associated
  - (**v**  $\Rightarrow$  in-memory, **i**  $\Rightarrow$  not-in-memory)
- Initially valid–invalid bit is set to **i** on all entries
- Example of a page table snapshot:



Frame #	valid-invalid bit
	<b>v</b>
	<b>v</b>
	<b>v</b>
	<b>v</b>
	<b>i</b>
....	
	<b>i</b>
	<b>i</b>

page table

- During address translation, if valid–invalid bit in page table entry is **i**  $\Rightarrow$  page fault



# Page Table When Some Pages Are Not in Main Memory






# Page Fault

- If there is a reference to a page,
  - first reference to that page will trap to operating system: *page fault*
    1. Operating system looks at another table to decide:
      - Invalid reference  $\Rightarrow$  abort
      - Just not in memory
    2. Get empty frame
    3. Swap page into frame
    4. Update the page table
    5. Set validation bit = **v**
    6. Restart the instruction that caused the page fault


# Page Fault (Cont.)

- Restart instruction

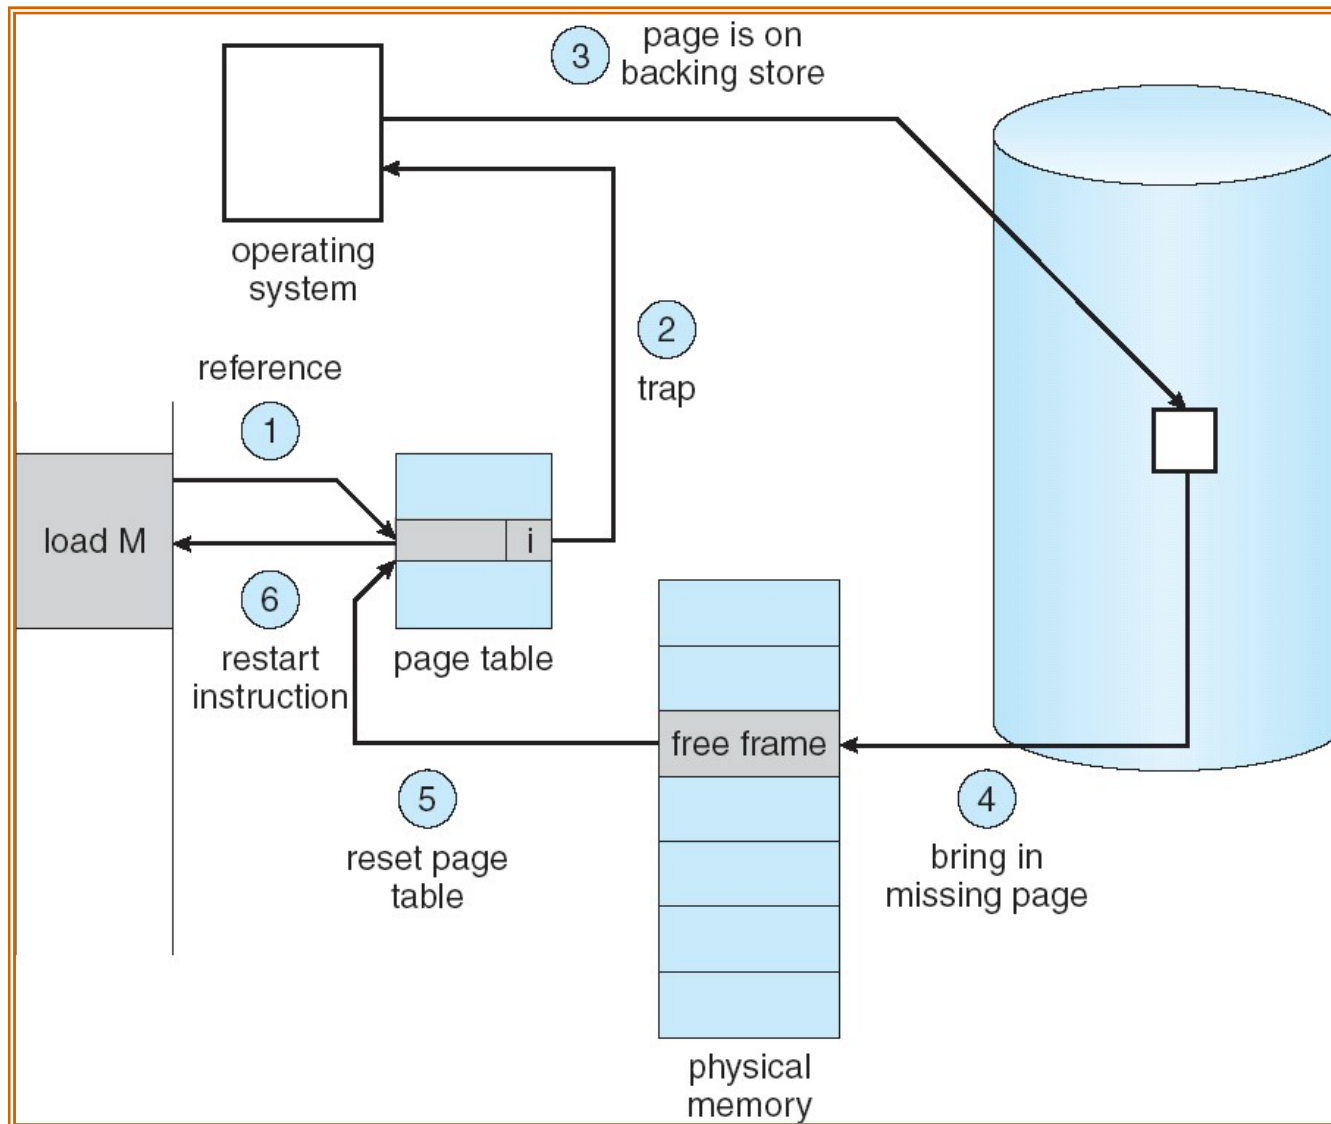
- block move



- auto increment/decrement location



# Steps in Handling a Page Fault



## Process information

Windows Task Manager

File Options View Shut Down Help


Applications Processes Performance Networking Users

Image Name	User Name	CPU	Mem Usage	Page Faults	VM Size
chrome.exe	t-thanh	08	90,672 K	438,921	86,972 K
Com4QLBEx.exe	SYSTEM	00	2,836 K	719	848 K
wmiprvse.exe	SYSTEM	00	5,108 K	2,152	1,912 K
chrome.exe	t-thanh	00	9,124 K	4,090	5,100 K
hpqwmie.exe	SYSTEM	00	4,008 K	1,032	2,244 K
iPodService.exe	SYSTEM	00	3,944 K	1,036	2,512 K
chrome.exe	t-thanh	00	18,432 K	7,143	11,900 K
TOTALCMD.EXE	t-thanh	00	13,480 K	7,506	8,524 K
chrome.exe	t-thanh	01	28,636 K	232,654	21,172 K
chrome.exe	t-thanh	00	16,088 K	7,297	9,872 K
chrome.exe	t-thanh	00	42,048 K	308,297	23,096 K
chrome.exe	t-thanh	01	40,264 K	417,093	30,652 K
BTTTray.exe	t-thanh	00	5,068 K	1,334	3,288 K
NMBgMonitor.exe	t-thanh	00	4,880 K	1,809	2,076 K
chrome.exe	t-thanh	00	38,796 K	49,788	31,804 K
smax4pnp.exe	t-thanh	00	4,244 K	1,362	2,460 K
iTunesHelper.exe	t-thanh	00	8,600 K	2,243	5,264 K
jusched.exe	t-thanh	00	2,432 K	619	824 K
SynTPEnh.exe	t-thanh	00	4,628 K	1,300	1,440 K
igfxsvc.exe	t-thanh	00	3,580 K	914	1,532 K
chrome.exe	t-thanh	00	66,256 K	818,886	46,860 K
openvpn-gui.exe	admin	00	3,948 K	1,014	1,260 K
openvpn.exe	admin	00	4,380 K	5,280	1,368 K
QLBCTRL.exe	t-thanh	00	5,196 K	1,670	2,592 K
igfxpers.exe	t-thanh	00	2,880 K	797	756 K
hkcmd.exe	t-thanh	00	3,392 K	880	1,004 K
igfxtray.exe	t-thanh	00	3,404 K	892	1,012 K
avgnt.exe	t-thanh	00	2,432 K	231,452	4,828 K
explorer.exe	t-thanh	01	24,468 K	21,466	15,892 K
wscntfy.exe	t-thanh	00	2,224 K	576	632 K
UniKeyNT.exe	t-thanh	00	3,424 K	884	1,140 K
chrome.exe	t-thanh	00	43,184 K	138,701	37,024 K
btwdins.exe	SYSTEM	00	2,356 K	939	1,940 K
svchost.exe	SYSTEM	00	24,368 K	26,110	15,180 K


☐ Show processes from all users

End Process

Processes: 62 CPU Usage: 14% Commit Charge: 1510M / 4948M



# Page fault in Windows



Task Manager

Type a name, publisher, or PID to search

Details

Run new task End task


Name	PID	Status	User name	CPU	Memory (a...	Page faults	Archite...	Description
MsMpEng.exe	3012	Running	SYSTEM	00	150,752 K	20		Number of page faults generated by the process since it was started
POWERPNT.EXE	5940	Running	thanhnt	00	88,476 K	1		
chrome.exe	2064	Running	thanhnt	00	83,212 K	102,215	x64	Google C...
dwm.exe	1188	Running	DWM-1	00	56,972 K	89,831	x64	Desktop ...
explorer.exe	6384	Running	thanhnt	01	29,224 K	57,615	x64	Windows ...
SearchHost.exe	6676	Suspended	thanhnt	00	0 K	48,494	x64	SearchHost
chrome.exe	7684	Running	thanhnt	00	22,208 K	45,386	x64	Google C...
Taskmgr.exe	9536	Running	thanhnt	02	37,084 K	43,371	x64	Task Mana...
svchost.exe	4812	Running	SYSTEM	00	11,780 K	42,070	x64	Host Proc...
svchost.exe	3728	Running	LOCAL SE...	00	4,764 K	40,480	x64	Host Proc...
vmware-hostd.exe	2968	Running	SYSTEM	00	24,164 K	34,226	x86	vmware-h...
Registry	104	Running	SYSTEM	00	5,720 K	34,004	x64	NT Kernel ...
csrss.exe	692	Running	SYSTEM	00	1,264 K	30,377	x64	Client Ser...
SystemSettings.exe	10224	Suspended	thanhnt	00	0 K	28,092	x64	SystemSet...
MoUsocoreWorker.e...	1352	Running	SYSTEM	00	4,364 K	27,468	x64	MoUsocore...
ShellExperienceHost...	1512	Suspended	thanhnt	00	0 K	27,224	x64	Windows ...
OneDrive.exe	7560	Running	thanhnt	00	15,456 K	24,811	x64	Microsoft ...
StartMenuExperienc...	6736	Running	thanhnt	00	24,504 K	24,089	x64	Windows ...
svchost.exe	2900	Running	LOCAL SE...	00	11,496 K	20,141	x64	Host Proc...
vmware-authd.exe	4000	Running	SYSTEM	00	2,344 K	18,733	x86	VMware A...
svchost.exe	3708	Running	SYSTEM	00	11,084 K	16,724	x64	Host Proc...
svchost.exe	2028	Running	SYSTEM	00	1,260 K	16,260	x64	Host Proc...
chrome.exe	8284	Running	thanhnt	00	10,632 K	16,245	x64	Google C...
svchost.exe	2540	Running	SYSTEM	00	7,820 K	16,141	x64	Host Proc...

# Question



- Which of the following is incorrect about a page fault?
  - A. it happens in paging on demand
  - B. it happens when a reference to a page that is not in MEM
  - C. when a page fault occurs the corresponding process will be terminated
  - D. a page fault handler is called whenever it occurs

## Question



- Which of the following is **incorrect order** of steps in handling a page fault?
  - A. check if the valid bit is invalid  $\Rightarrow$  raise a page fault
  - B. a page fault is raised  $\Rightarrow$  find the page in backing store
  - C. a page fault is raised  $\Rightarrow$  find a free frame
  - D. load the page into memory  $\Rightarrow$  restart the instruction





# If no free frame available

- Call page replacement procedure
  - swap out an **unused** page from MEM
- Algorithms
  - Optimal, FIFO, LRU, LRU-approximation, counting
- Performance of the algorithm
  - page-fault rate
  - which algorithm is better?

# Performance of paging on demand



- Page Fault Rate  $0 \leq p \leq 1.0$ 
  - if  $p = 0$  no page faults
  - if  $p = 1$ , every reference is a page fault
- Effective Access Time (EAT)
$$\text{EAT} = (1 - p) * \text{memory access} \\ + p (\text{page fault overhead} \\ + \text{swap page out} \\ + \text{swap page in} \\ + \text{restart overhead})$$

# Paging on Demand Example




- Memory access time = 200 nanoseconds
- Average page-fault service time = 8 milliseconds
  - $EAT = (1 - p) \times 200 + p (8 \text{ milliseconds})$   
 $= (1 - p) \times 200 + p \times 8,000,000$   
 $= 200 + p \times 7,999,800$
- If one access out of 1,000 causes a page fault
  - $EAT = 8.2 \text{ microseconds.}$
  - slowdown by a factor of 41!!

# Page Replacement




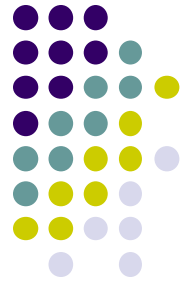
- Prevent over-allocation of memory
  - include page replacement in page-fault service routine
- Use **modify (dirty) bit** to reduce overhead of page transfers
  - only **modified** pages are written to disk
- Page replacement completes separation between logical memory and physical memory
  - large virtual memory can be provided on a smaller physical memory

# Page Replacement



```
ntthanh@turing:~  
HighTotal:      0 kB  
HighFree:       0 kB  
LowTotal:      1542448 kB  
LowFree:       299084 kB  
SwapTotal:     3112952 kB  
SwapFree:     3112952 kB  
Dirty:         24 kB  
Writeback:      0 kB  
AnonPages:     180092 kB  
Mapped:        29504 kB  
Slab:          90524 kB  
PageTables:     6912 kB  
NFS_Unstable:   0 kB  
Bounce:         0 kB  
CommitLimit:   3884176 kB  
Committed_AS:  438656 kB  
VmallocTotal: 34359738367 kB  
VmallocUsed:    2220 kB  
VmallocChunk: 34359735923 kB  
HugePages_Total: 0  
HugePages_Free: 0  
HugePages_Rsvd: 0  
Hugepagesize:  2048 kB  
[ntthanh@turing ~]$
```


# Need For Page Replacement




# Basic Page Replacement



1. Find the location of the desired page on disk
2. Find a free frame
  - If there is a free frame, use it
  - Else use a page replacement algorithm to select a **victim** frame (and swap it out)
3. Bring the desired page into the (newly) free frame; update the page table
4. Resume the process



# Page Replacement





# Question



Which of the following is **incorrect** about page replacement?

- A. a victim frame is selected to be swapped out
- B. the page table which is the victim will be updated
- C. the victim frame is always written into the backing store
- D. the victim frame is only written into the backing store if it is dirty




# Page Replacement Algorithms

- Want lowest page-fault rate
- Evaluate algorithm
  - run it on a particular string of memory references (reference string)
  - compute the number of page faults on that string
- In all our examples, the reference string is

**1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5**

# Graph of Page Faults Versus The Number of Frames





# Optimal Algorithm

- Replace page that **will not** be used for longest period of time
- 4 frames example

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

How do you know this?

- Used for measuring **how well** the algorithm performs

1
2
3
4

4

5

6 page faults



## Question

- A reference string 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1, Optimal algorithm is used with 3 frames. Which of the following is the correct order of swapped out pages?
  - A. 7 1 0 3 4 2
  - B. 7 0 1 4 3 2
  - C. 7 1 0 4 3 2
  - D. 7 1 4 1 3 2



## Question

- A reference string 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1, Optimal algorithm is used with 3 frames. Which of the following is the correct number of page faults?
  - A. 8
  - B. 9
  - C. 10
  - D. 11

# First-In-First-Out (FIFO) Algorithm



- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

- 3 frames

1	1	4	5
2	2	1	3
3	3	2	4

9 page faults

- 4 frames


1	1	5	4
2	2	1	5
3	3	2	
4	4	3	

10 page faults  $\Rightarrow$  Belady's anomaly

Belady's Anomaly: more frames  $\Rightarrow$  more page faults




# FIFO Illustrating Belady's Anomaly





# FIFO Page Replacement





## Question

- A reference string 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1, FIFO is used with 3 frames. Which of the following is the correct order of swapped out pages?
  - A. 7 0 1 2 3 0 4 2 3 0 1 2
  - B. 7 0 1 2 3 0 4 3 2 0 1 2
  - C. 7 0 1 3 2 0 4 2 3 0 1 2
  - D. 7 0 1 2 3 0 4 2 3 1 0 2



## Question

- A reference string 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1, FIFO is used with 3 frames. Which of the following is the correct number of page faults?
  - A. 13
  - B. 14
  - C. 15
  - D. 16

# Least Recently Used (LRU) Algorithm



- Least recently used page is swapped out first
  - Reference string: 1, 2, 3, 4, 1, 2, **5**, 1, 2, **3**, **4**, **5**
  - 4 frames

1	1	1	1	<b>5</b>
2	2	2	2	2
3	<b>5</b>	5	<b>4</b>	4
4	4	<b>3</b>	3	3



## Question

- A reference string 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1, LRU is used with 3 frames.
- Which of the following is the correct order of swapped out pages?
  - A. 7 1 2 3 1 4 1 3 2
  - B. 7 2 1 3 0 4 2 3 2
  - C. 7 1 2 3 0 4 1 2 3
  - D. 7 1 2 3 0 4 0 3 2



## Question

- A reference string 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1, LRU is used with 3 frames.
- Which of the following is the correct number of page faults?
  - A. 13
  - B. 12
  - C. 11
  - D. 10



# Least Recently Used Algorithm




- Counter implementation
  - Every page entry has a counter;
    - every time page is referenced, copy the clock into the counter
  - When a page needs to be swapped
    - look at the counters to determine



## LRU Algorithm (Cont.)

- Stack implementation
  - keep a stack of page numbers in a double link form
  - Page referenced:
    - move it to the top
    - requires 6 pointers to be changed
  - No search for replacement

# Use Of A Stack to Record The Most Recent Page References

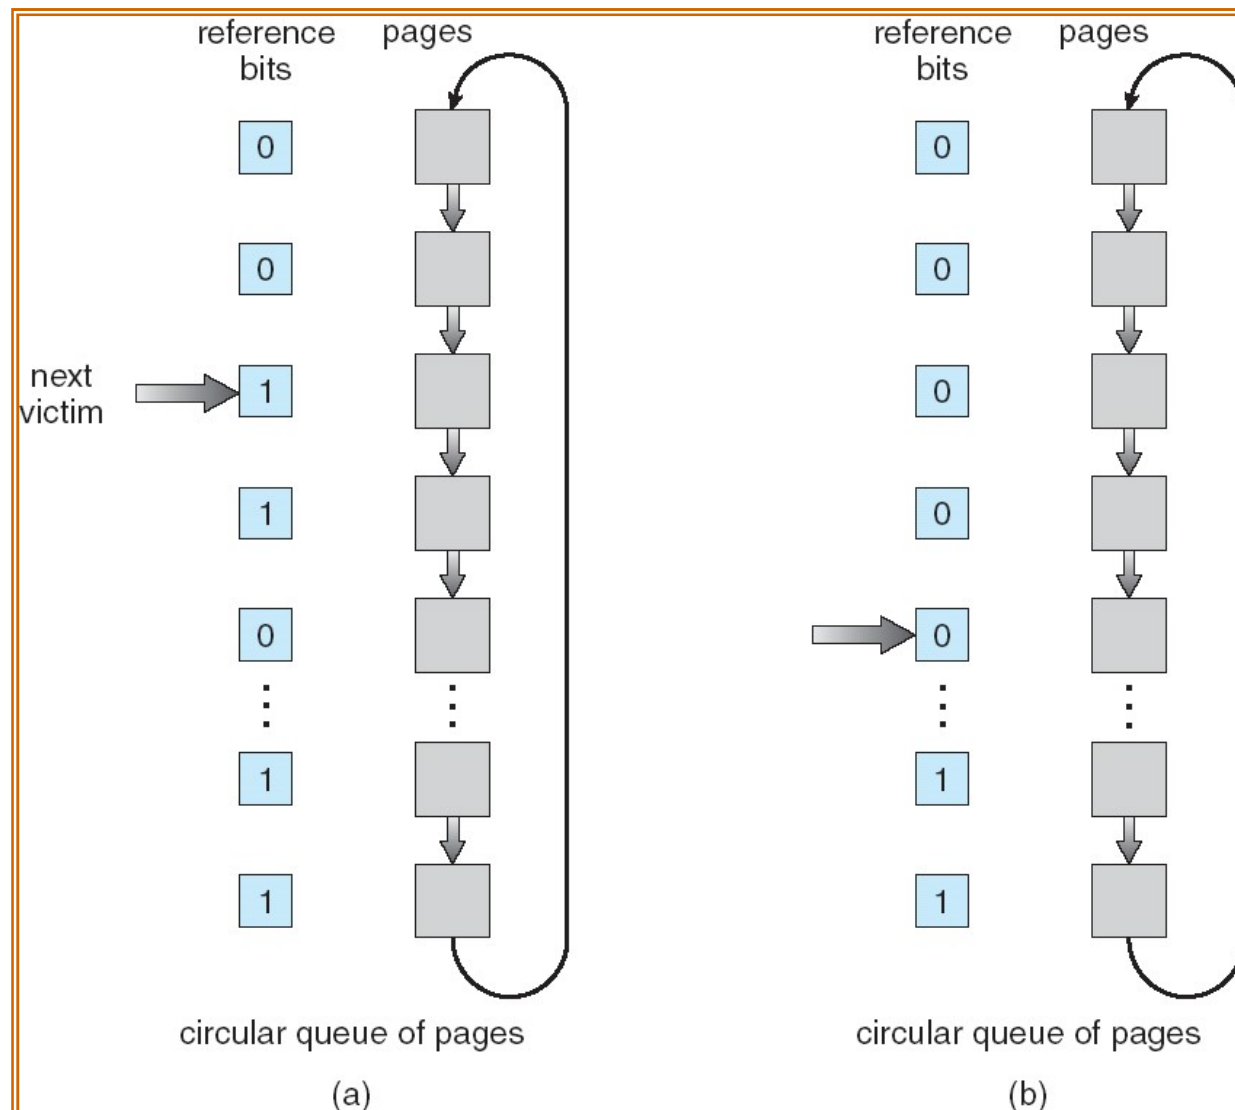


# LRU Approximation Algorithms



- Reference bit
  - With each page associate a bit, initially = 0
  - When page is referenced bit set to 1
  - Replace the one which is 0 (if one exists)
    - We do not know the order
- Second chance (follow clock order)
  - Need reference bit
  - Clock replacement
  - If page at the pointer has reference bit = 1 then:
    - set reference bit 0
    - leave page in memory
    - Move the pointer to next page, subject to same rules
  - Else select the page as the victim

# Second-Chance (clock) Page-Replacement Algorithm





## Question

- Suppose the second chance is used;
  - the reference bits of frames are: 1 1 0 1 1 0
  - the head is at second frame
- Which of the following are the reference bits after a page replacement is done
  - A. 0 0 0 1 1 0
  - B. 1 0 1 1 1 0
  - C. 1 0 0 1 1 0
  - D. 1 0 1 0 1 0

# Question



- Suppose the LRU counter implementation without additional support, which of the following is incorrect?
  - A. Timestamp is used to mark the referred time
  - B. The smallest timestamp is selected for replacement
  - C. Searching is need to find the smallest timestamp
  - D. No need to search to find the smallest timestamp



# Question



- Suppose the LRU counter implementation without additional support, which of the following is incorrect?
  - A. Timestamp is used to mark the referred time
  - B. The smallest timestamp is selected for replacement
  - C. Searching is need to find the smallest timestamp
  - D. No need to search to find the smallest timestamp

# Counting Algorithms



- Keep a counter of the number of references that have been made to each page
- **Least Frequently Used (LFU) Algorithm**
  - replaces page with smallest count
- **Most Frequently Used (MFU) Algorithm**
  - based on the argument that the page with the smallest count was probably just brought in and has yet to be used

# Allocation of Frames



- Each process needs *minimum* number of pages
- Example: IBM 370 – 6 pages to handle SS MOVE instruction:
  - instruction is 6 bytes, might span 2 pages
  - 2 pages to handle *from*
  - 2 pages to handle *to*
- Two major allocation schemes
  - fixed allocation
  - priority allocation

# Fixed Allocation



- Equal allocation
  - For example, if there are 100 frames and 5 processes, give each process 20 frames.
- Proportional allocation
  - Allocate according to the size of process

- $s_i$  = size of process  $p_i$
- $S = \sum s_i$
- $m$  = total number of frames
- $a_i$  = allocation for  $a_i = \frac{s_i}{S} \times m$

$$m = 64$$

$$s_1 = 10$$

$$s_2 = 127$$

$$a_1 = \frac{10}{137} \times 64 \approx 5$$

$$a_2 = \frac{127}{137} \times 64 \approx 59$$

# Question



- A system uses proportional allocation and has
  - 90 frames x 2KB
  - 3 processes with size of (138KB, 96KB, 164KB)
- Which of the following is the correct number of allocated frames of ( $P_1$ ,  $P_2$ ,  $P_3$ )
  - A. 32, 21, 37
  - B. 31, 22, 37
  - C. 30, 22, 38
  - D. 33, 22, 35

# Priority Allocation



- Use a proportional allocation scheme using priorities rather than size
- If process  $P_i$  generates a page fault,
  - select for replacement one of its frames
  - select for replacement a frame from a process with **lower priority** number

# Global vs. Local Allocation



- **Global replacement**

- process selects a replacement frame from the set of all frames;
- one process can take a frame from another

- **Local replacement**

- each process selects from only its own set of allocated frames




# Thrashing



- If a process does not have “enough” frames, the **page-fault rate is very high**.
- **Thrashing**
  - a process is **busy swapping** pages in and out
- This caused by:
  - low CPU utilization
  - operating system thinks that it needs to increase the degree of multiprogramming
  - another process added to the system

# Thrashing (Cont.)



# Solutions to Thrashing



- Use local allocation
- Use priority allocation
  - not good solution
- Working set model
  - A suitable solution



## Question

- Which of the following is incorrect about priority allocation?
  - A. higher priority process is allocated first
  - B. it prevents thrashing from happening
  - C. frames are allocated globally
  - D. it does not prevent thrashing from happening




# Working-Set Model

- $\Delta \equiv$  working-set window
  - a number of page references, e.g. 10,000
- Working set of Process  $P_i$ 
  - $WSS_i$  = total number of pages referenced in the most recent  $\Delta$  (varies in time)
  - if  $\Delta$  too small will not encompass entire locality
  - if  $\Delta$  too large will encompass several localities
  - if  $\Delta = \infty \Rightarrow$  will encompass entire program
- $D = \sum WSS_i \equiv$  total demand frames
- if  $D > m$  (total of frames)  $\Rightarrow$  Thrashing
- Policy if  $D > m$ , then suspend one process



# Working-set model



# Keep Track of the WSet



- Approximate with interval timer + a reference bit
- Example:  $\Delta = 10,000$ 
  - Timer interrupts after every 5000 time units
  - Keep in memory 2 bits for each page in last 10000 and 15000
  - 15000---..bit1..-- 10000 ---..bit2..--- 5000 ---..ref.. --- now
  - Whenever a timer interrupts copy and sets the values of all reference bits to 0
  - If one of the bits in memory = 1  $\Rightarrow$  page in working set
- Why is this not completely accurate?
- Improvement = 10 bits and interrupt every 1000 time units

# Question



- Suppose a delta = 10; reference string
  - 2 6 1 5 7 7 7 7 5 1 6 2 3 4 4 4 3 4 4 4 1 3 2 3 4 4 4 3  
4 4 4 ...
- Which of the following is the correct WSS at 20<sup>th</sup> reference?
  - A. {2 3 4 6}
  - B. {2 3 4 5 6}
  - C. {1 2 3 4 6}
  - D. {7 1 2 3 4 6}





Question?