


# Operating System

---


**Nguyen Tri Thanh**  
**ntthanh@vnu.edu.vn**



# File System Implementation

---

File-System Structure  
File-System Implementation  
Directory Implementation  
Allocation Methods  
Free-Space Management  
Efficiency and Performance  
Recovery





# Objectives

- Introduce what a file is
- Introduce file system implementation
- Introduce directory implementation
- Introduce 3 allocation methods
- Introduce free space management

# Reference



- Chapter 10, 11 of **Operating System Concepts**

# File Concept



- Contiguous logical address space
- Types:
  - Data
    - numeric
    - character
    - binary
  - Program

# File Structure



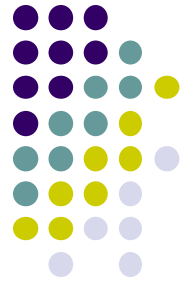
- None - sequence of words, bytes
- Simple record structure
  - Lines
  - Fixed length
  - Variable length
- Complex Structures
  - Formatted document
  - Relocatable load file
- Can simulate last two with first method by inserting appropriate control characters
- Who decides:
  - Operating system
  - Program

# File Attributes



- Name
  - only information kept in human-readable form
- Identifier
  - unique (number) identifies file within file system
- Type
  - needed for systems that support different types
- Location
  - pointer to file location on storage device

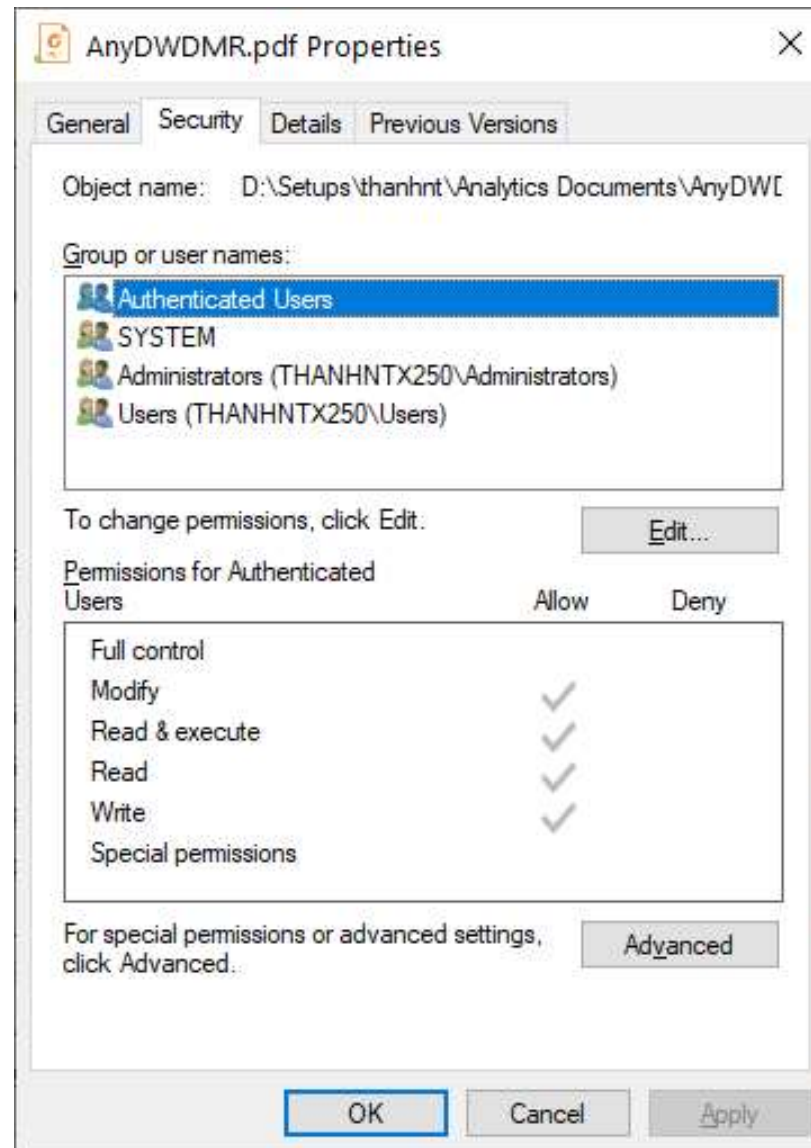
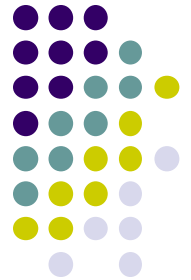
# File Attributes (cont'd)



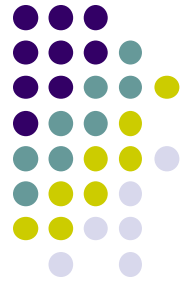
- Size
  - current file size
- Protection
  - controls who can do reading, writing, executing
- Time, date, and user identification
  - data for protection, security, and usage monitoring
- Information about files are kept in the **directory structure** on the disk



# File Attributes (cont'd)



# File System Structure



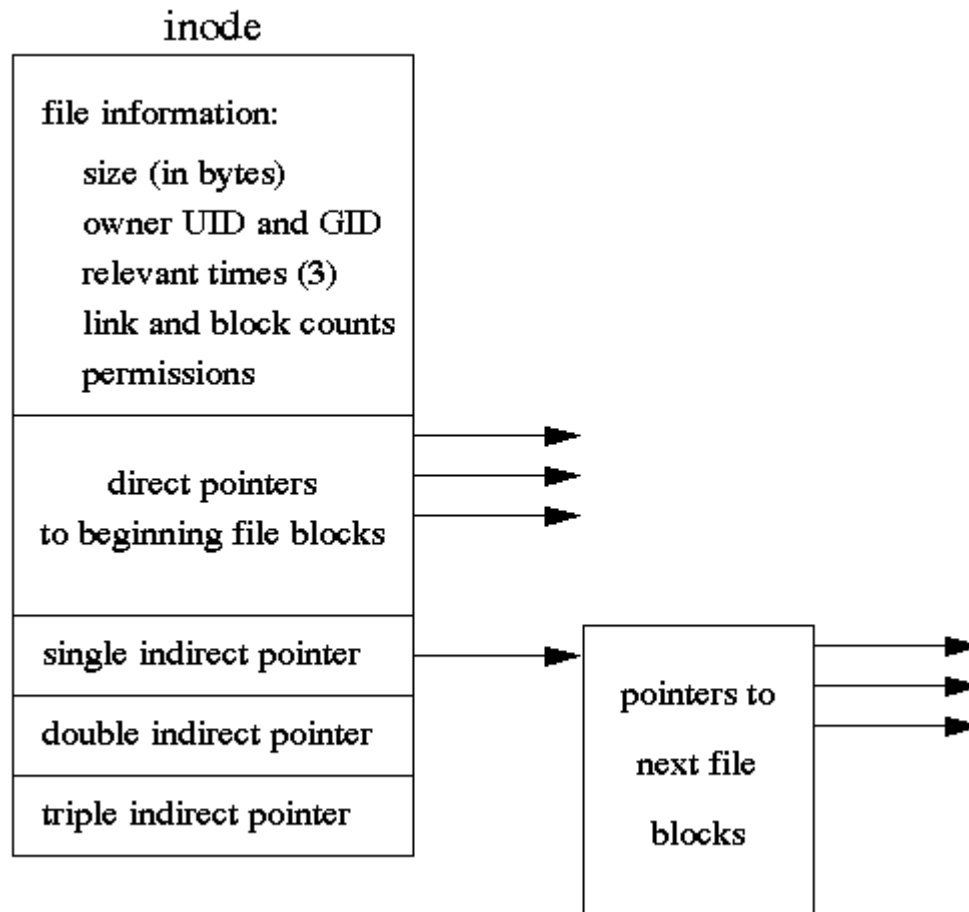
- File system resides on secondary storage
  - HDD disks, CD ROM, DVD, flash drive, SSD, ...
- File system organized into layers
- File control block (FCB)
  - storage structure of information about a file
  - *inode* (index node) has partial information of FCB on Linux



# A Typical File Control Block

file permissions
file dates (create, access, write)
file owner, group, ACL
file size
file data blocks or pointers to file data blocks

# Inode on Linux

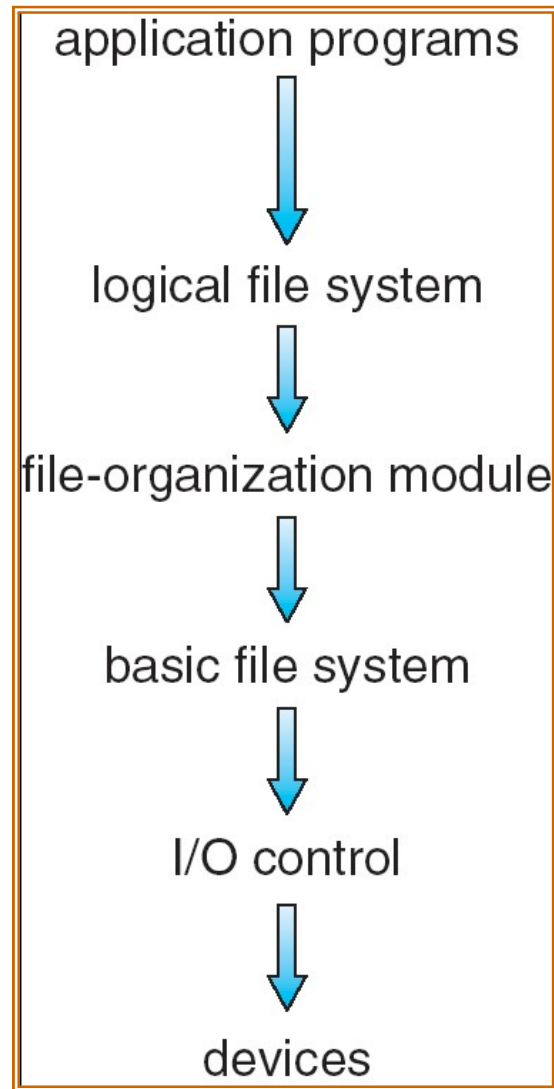




## Question

- Which of the following is incorrect about file control block (FCB)?
  - A. a data structure containing information of a file
  - B. OS needs a FCB to access a file
  - C. FCB of a file is updated when a file is accessed
  - D. FCBs reside in memory

# Layered File System

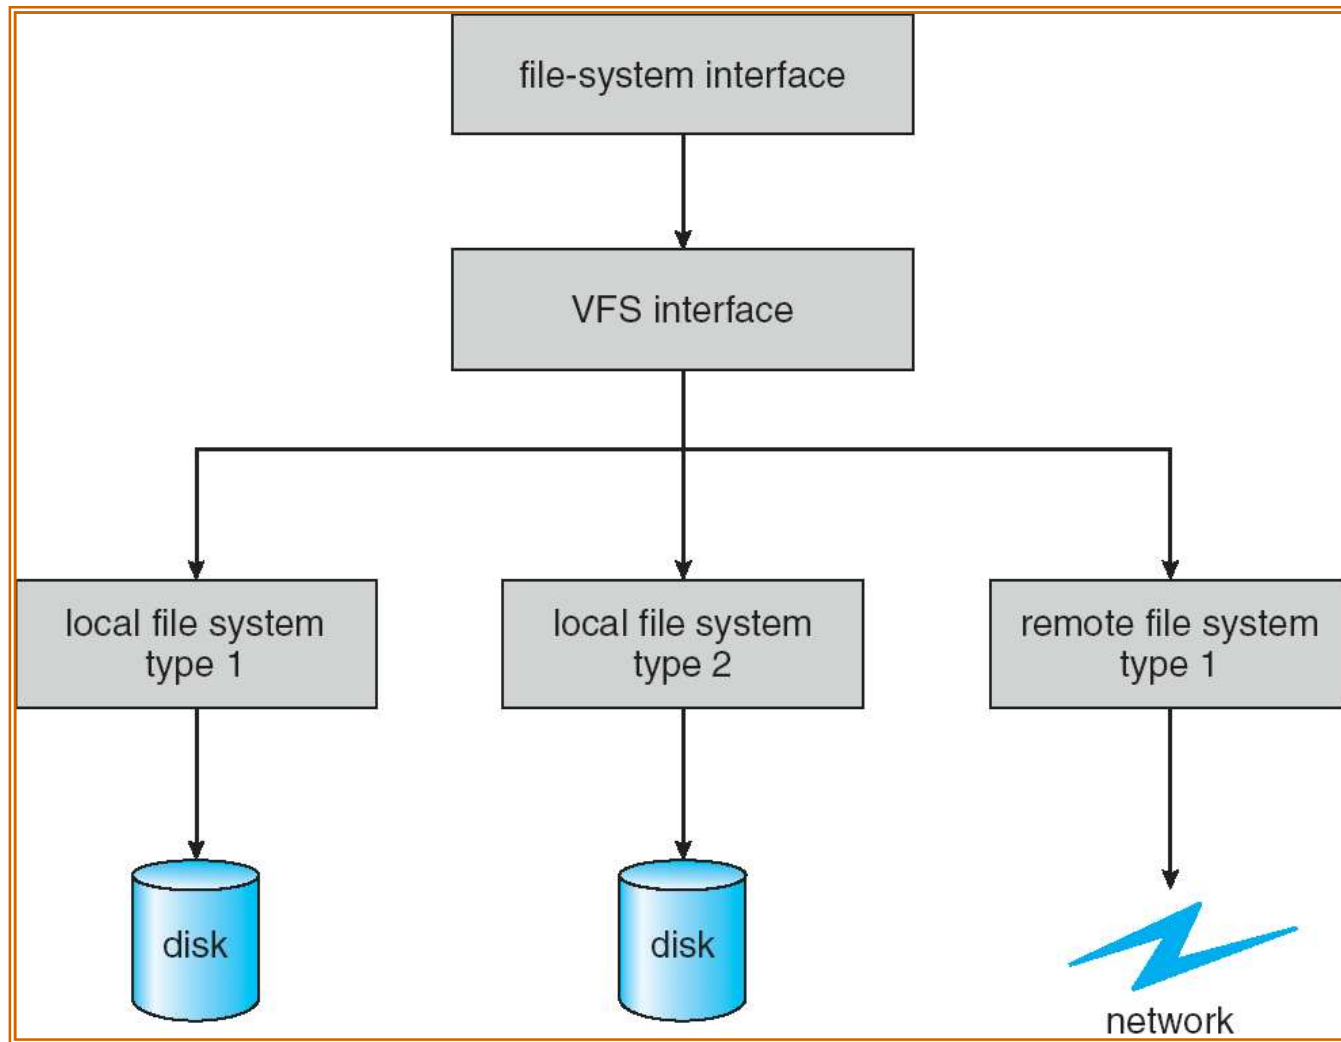



# Virtual File Systems



- Virtual File Systems (VFS)
  - provide an object-oriented way of implementing file systems
  - allow the same system call interface (the API) to be used for different types of file systems
  - the API is to the VFS interface, rather than any specific type of file system.

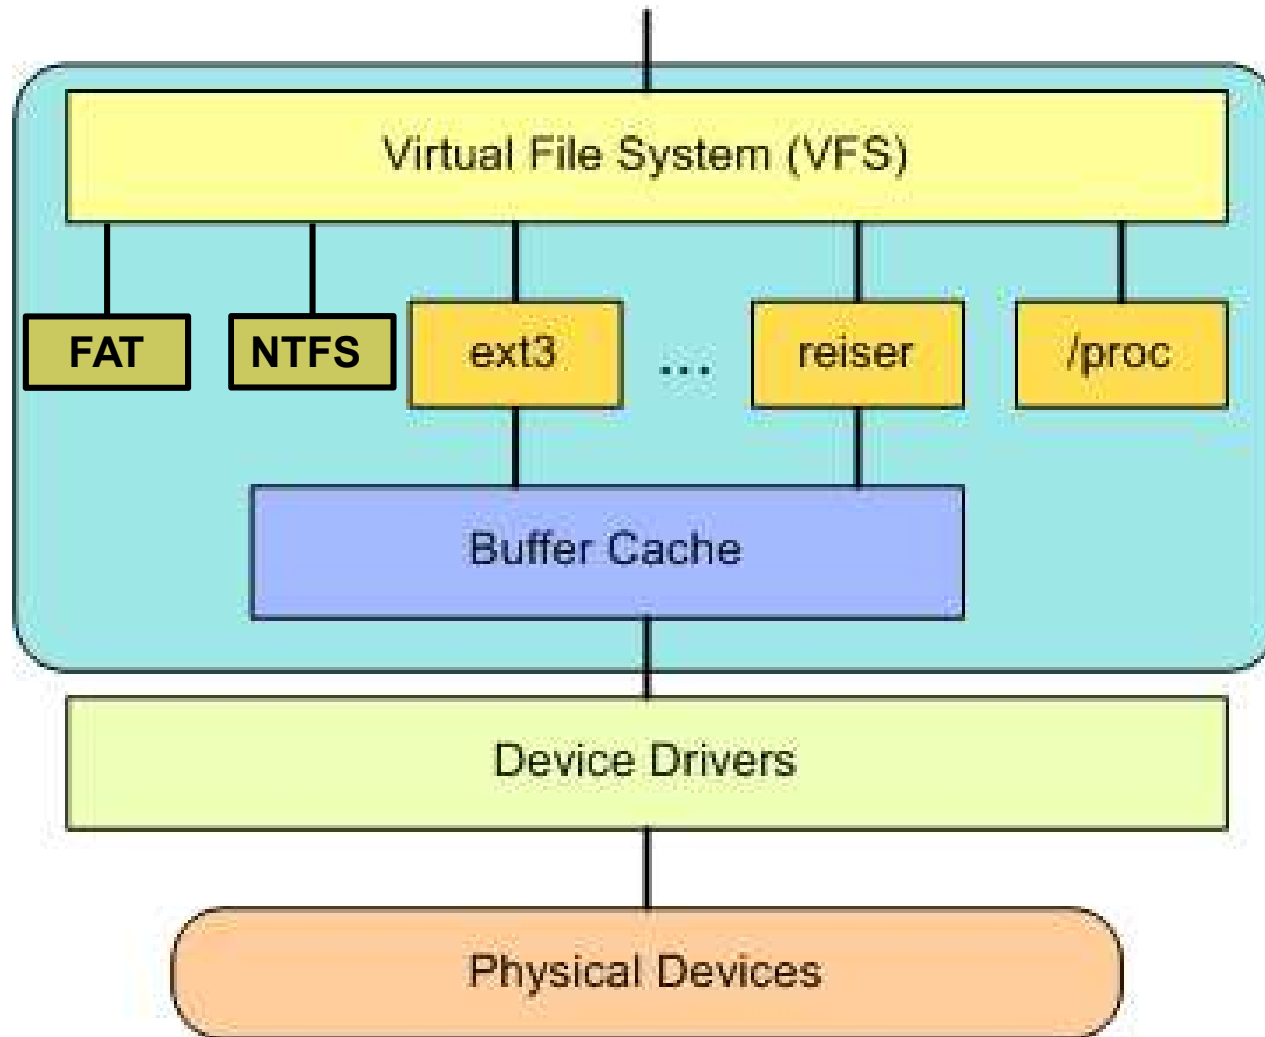
# Schematic View of Virtual File System







# Linux VFS



# File system list



W List of file systems - Wikipedia x + You are screen sharing Stop Share

en.wikipedia.org/wiki/List\_of\_file\_systems

Languages العربية Deutsch Español 한국어 Русский తెలుగు 中文 4 more Edit links

Disk file systems are usually block-oriented. Files in a block-oriented file system are sequences of blocks, often featuring fully random-access read, write, and modify operations.

- **ADFS** – Acorn's Advanced Disc filing system, successor to DFS.
- **AdvFS** – Advanced File System, designed by Digital Equipment Corporation for their Digital UNIX (now Tru64 UNIX) operating system.
- **APFS** – Apple File System is a next-generation file system for Apple products.
- **AthFS** – AtheOS File System, a 64-bit journaled filesystem now used by Syllable. Also called AFS.
- **BFS** – the Boot File System used on System V release 4.0 and UnixWare.
- **BFS** – the Be File System used on BeOS, occasionally misnamed as BeFS. Open source implementation called OpenBFS is used by the Haiku operating system.
- **Btrfs** – is a copy-on-write file system for Linux announced by Oracle in 2007 and published under the GNU General Public License (GPL).
- **CFS** – The Cluster File System from Veritas, a Symantec company. It is the parallel access version of VxFS.
- **CP/M file system** — Native filesystem used in the CP/M (Control Program for Microcomputers) operating system which was first released in 1974.
- **DOS 3.x** – Original floppy operating system and file system developed for the Apple II.
- **Extent File System (EFS)** – an older block filing system under IRIX.
- **ext** – Extended file system, designed for Linux systems.
- **ext2** – Second extended file system, designed for Linux systems.
- **ext3** – A journaled form of ext2.
- **ext4** – A follow up for ext3 and also a journaled filesystem with support for extents.
- **ext3cow** – A versioning file system form of ext3.
- **FAT** – File Allocation Table, initially used on DOS and Microsoft Windows and now widely used for portable USB storage and some other devices; **FAT12**, **FAT16** and **FAT32** for 12-, 16- and 32-bit table depths.
  - **VFAT** – Optional layer on Microsoft Windows FAT system to allow long (up to 255 character) filenames instead of only the 8.3 filenames allowed in the plain FAT filesystem.
  - **FATX** – A modified version of Microsoft Windows FAT system that is used on the original Xbox console.
- **FFS (Amiga)** – Fast File System, used on Amiga systems. This FS has evolved over time. Now counts FFS1, FFS Intl, FFS DCache, FFS2.

# File system list



W List of file systems - Wikipedia x + You are screen sharing Stop Share

en.wikipedia.org/wiki/List\_of\_file\_systems

- [ISO 9660](#) – Used on CD-ROM and DVD-ROM discs (rock-ridge and some are extensions to this)
- [JFS](#) – IBM Journaling file system, provided in Linux, OS/2, and AIX. Supports [extents](#).
- [LFS](#) – 4.4BSD implementation of a log-structured file system
- [MFS](#) – Macintosh File System, used on early Classic Mac OS systems. Succeeded by Hierarchical File System (HFS).
- [Next3](#) – A form of [ext3](#) with snapshots support.<sup>[6]</sup>
- [MFS](#) – TiVo's Media File System, a proprietary fault tolerant format used on TiVo hard drives for real time recording from live TV.
- [Minix file system](#) – Used on Minix systems
- [NILFS](#) – Linux implementation of a log-structured file system
- [NTFS](#) – (New Technology File System) Used on Microsoft's Windows NT-based operating systems
- [NetWare File System](#) – The original NetWare 2.x–5.x file system, used optionally by later versions.
- [NSS](#) – Novell Storage Services. This is a new 64-bit [journaling file system](#) using a balanced tree algorithm. Used in NetWare versions 5.0-up and recently ported to Linux.
- [OneFS](#) – One File System. This is a fully journaled, distributed file system used by [Isilon](#). OneFS uses FlexProtect and [Reed-Solomon](#) encodings to support up to four simultaneous disk failures.
- [OFS](#) – Old File System, on Amiga. Good for floppies, but fairly useless on hard drives.
- [OS-9 file system](#)
- [PFS](#) – and PFS2, PFS3, etc. Technically interesting file system available for the [Amiga](#), performs very well under a lot of circumstances. Very simple and elegant.
- [ProDOS](#) – Operating system and file system successor to [DOS 3.x](#), for use on Apple's computers prior to the Macintosh & Lisa computers, the Apple series, including the [IIgs](#)
- [Qnx4fs](#) – File system that is used in [QNX](#) version 4 and 6.
- [ReFS \(Resilient File System\)](#) – New file system by Microsoft that is built on the foundations of [NTFS](#) (but cannot boot, has a default cluster size of 64 KB and does not support compression) and is intended to be used with the [Windows Server 2012](#) operating system.
- [ReiserFS](#) – File system that uses [journaling](#)
- [Reiser4](#) – File system that uses [journaling](#), newest version of ReiserFS
- [Reliance](#) – Datalight's transactional file system for high reliability applications



## Question

- Which of the following is incorrect about VFS?
  - A. VFS allows an OS to support many different file systems
  - B. VFS provides the same API for all file systems
  - C. VFS is available in all OSes
  - D. VFS hides the detailed implementation of each file system from programmers




## Question

- Which of the following is a correct description of a directory?
  - A. a directory is a disk partition to contain files
  - B. a directory is actually a file containing partial information about its files
  - C. a directory is a container of files' data
  - D. a directory contains the FCB and data of files



# Directory Structure

- A collection of **nodes** containing information about all files



# Directory Implementation



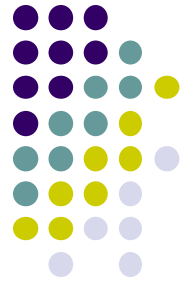
- **Linear list**

- list of file names with pointer to the data blocks
- simple to program
- time-consuming to execute
- FAT [http://en.wikipedia.org/wiki/File\\_Allocation\\_Table](http://en.wikipedia.org/wiki/File_Allocation_Table)
- Linux FS (Ext3)

- **Hash Table**

- linear list with hash data structure
- decreases directory search time
- **collisions** resolution
- fixed size

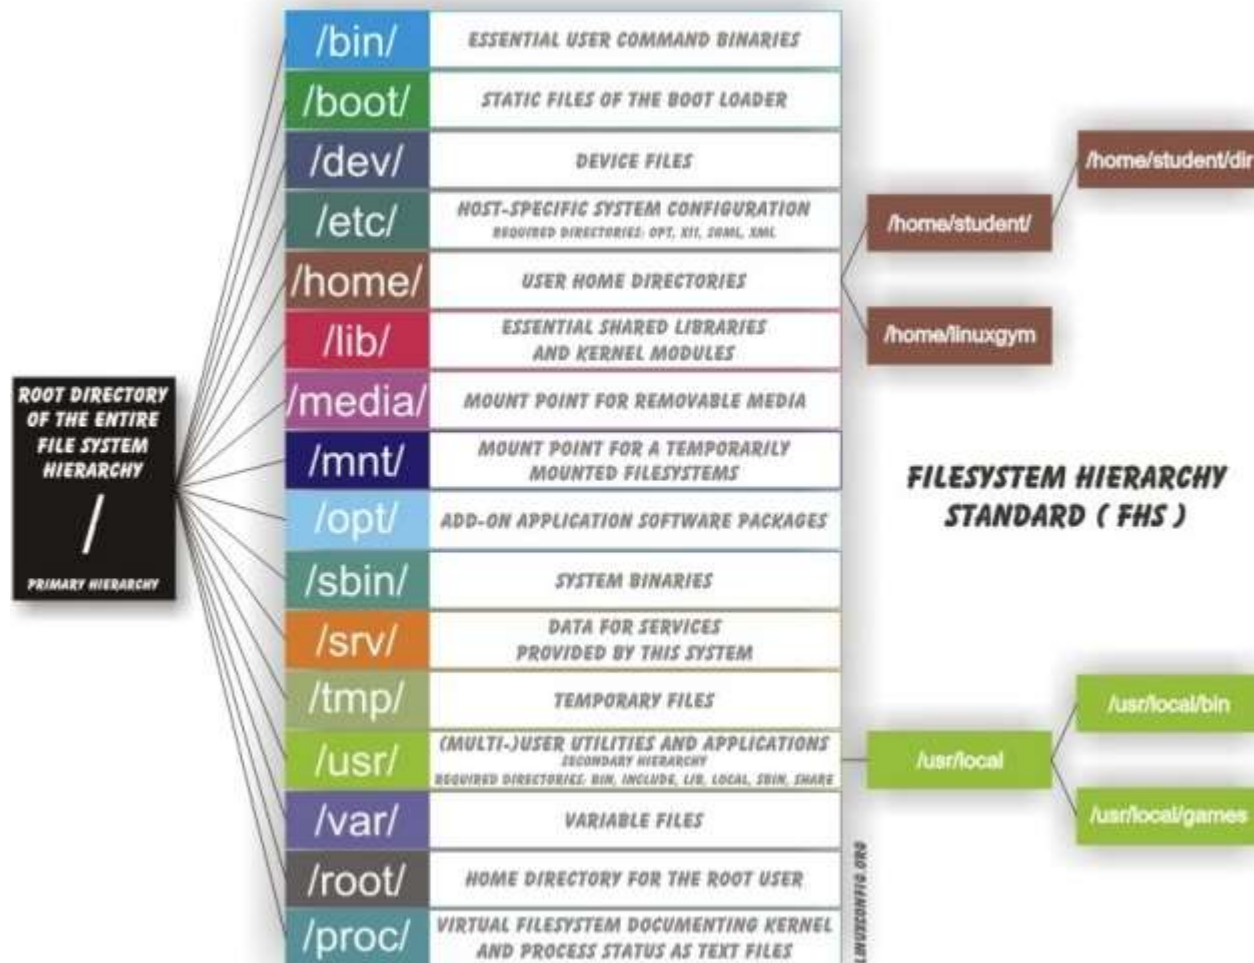

# Directory Implementation



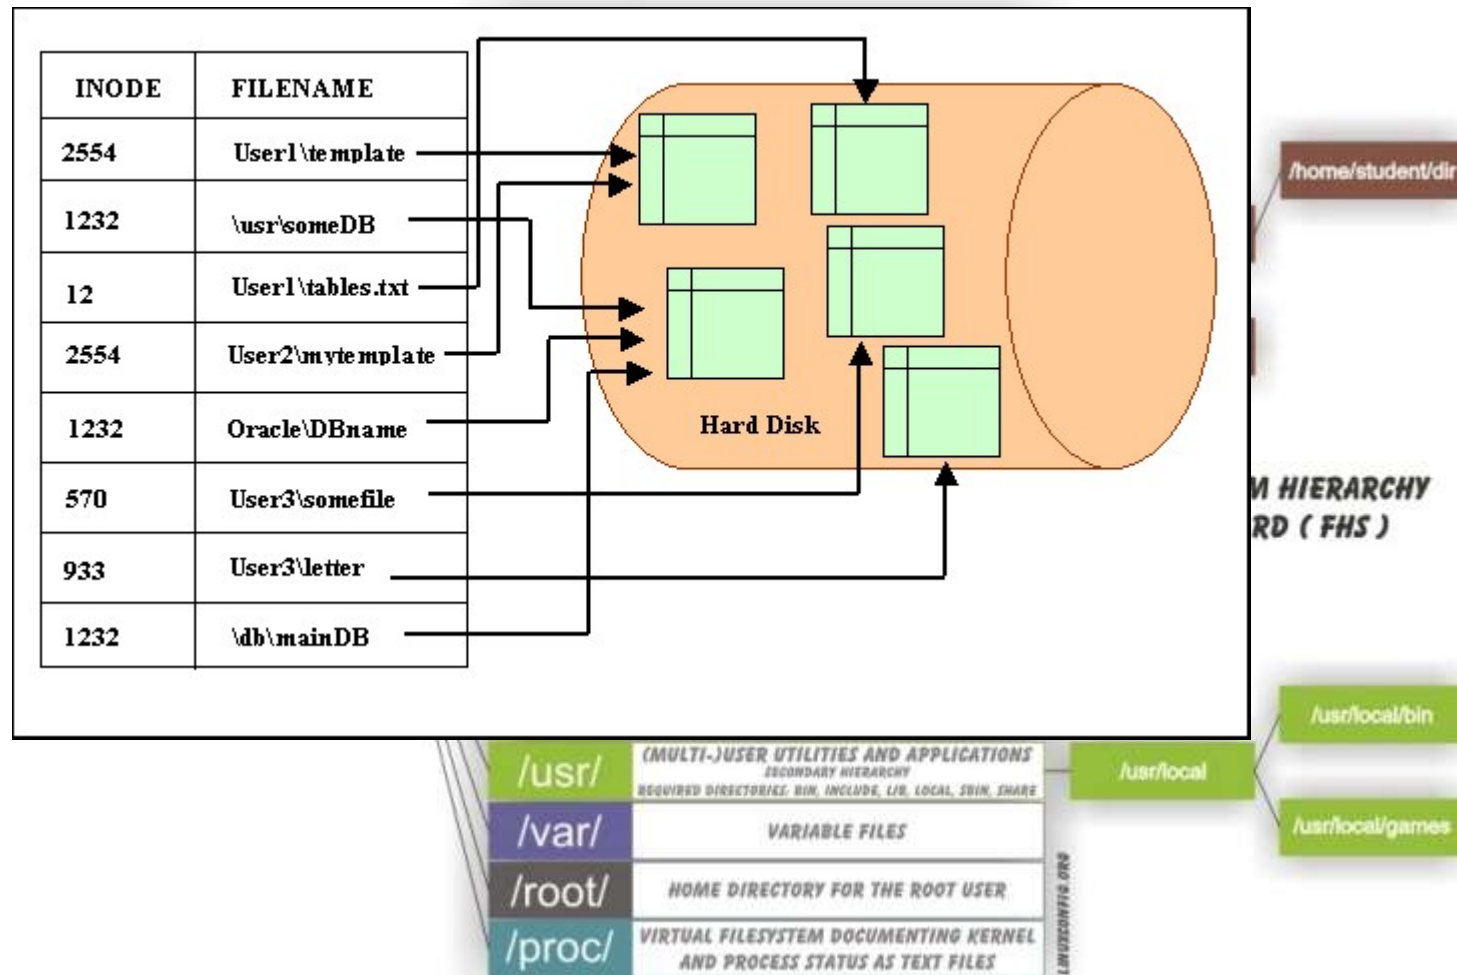
- **Balanced binary tree**
  - RAISERFS <http://en.wikipedia.org/wiki/ReiserFS>



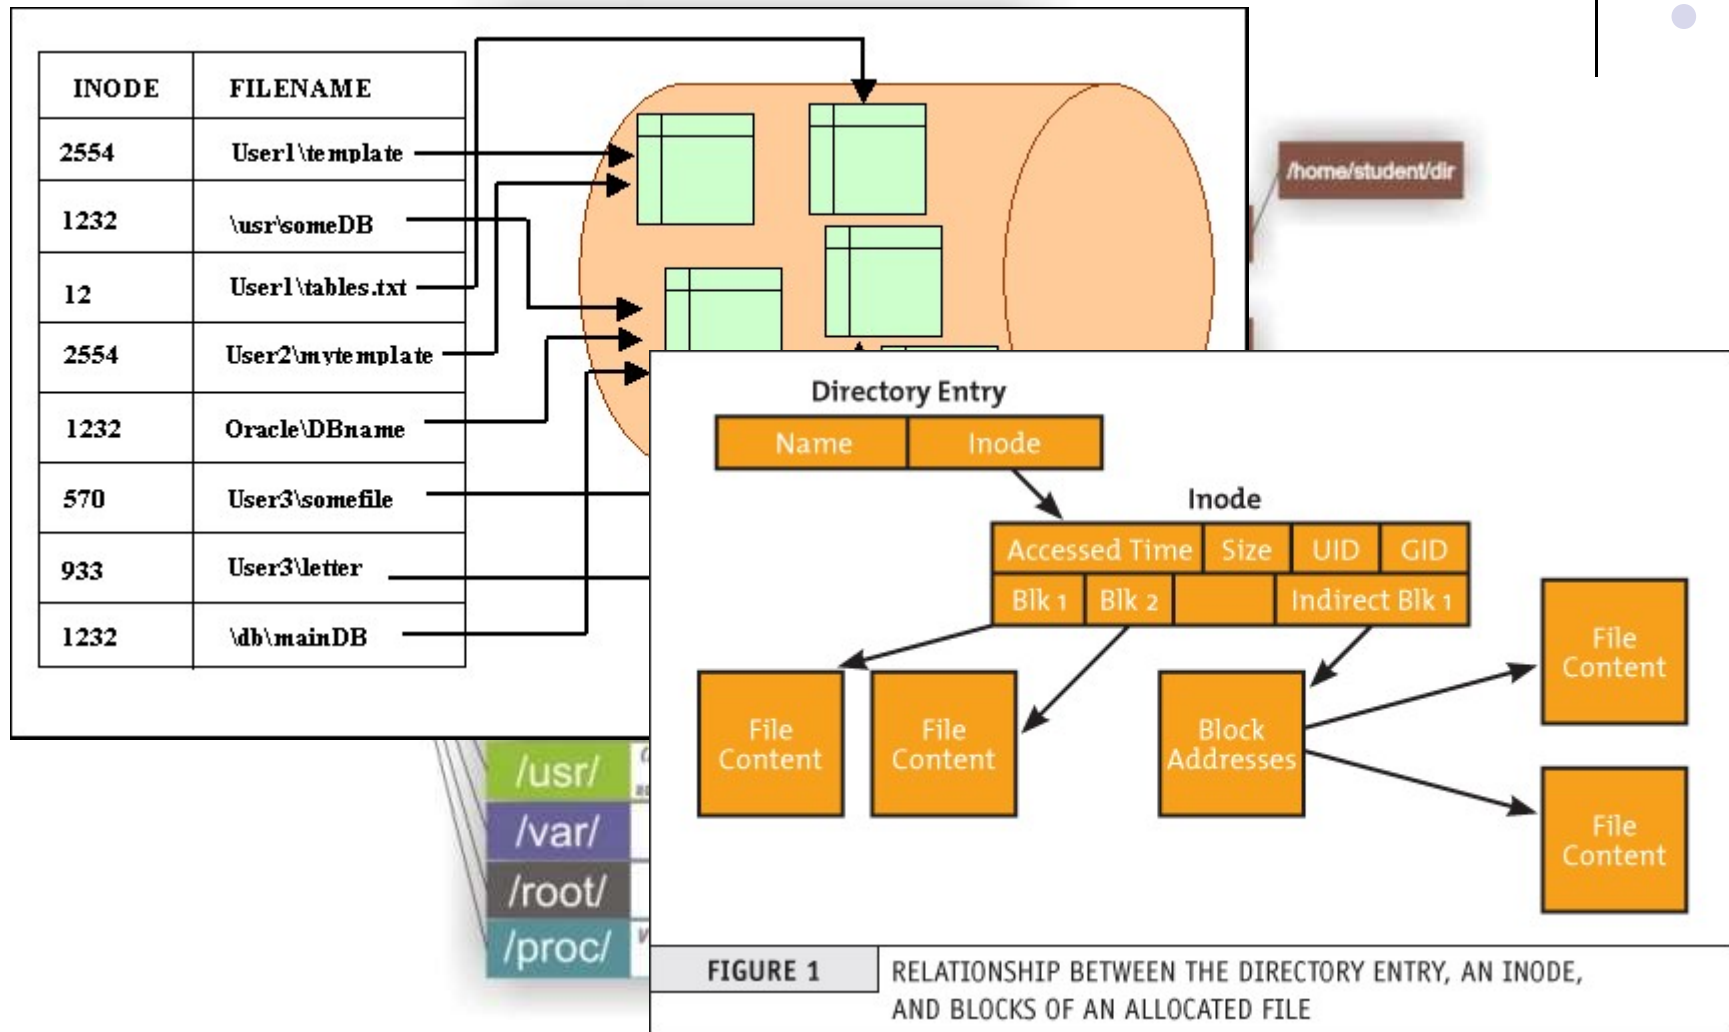
# Directory in Linux



# Directory in Linux



# Directory in Linux





# **File management API**

# File Operations



- Create
- Write
- Read
- Reposition within file
- Delete
- Truncate
- $Open(F_i)$ 
  - search the directory structure on disk for entry  $F_i$ , and move the content of entry to memory
- $Close(F_i)$ 
  - move the content of entry  $F_i$  in memory to directory structure on disk

# File Operations



File Management Functions - Win x +

docs.microsoft.com/en-us/windows/win32/fileio/file-management-functions

Filter by title

file management functions

- File Management Functions
- AddUsersToEncryptedFile
- AreFileApisANSI
- Canceled
- CanceledEx
- CancelSynchronousIo
- CheckNameLegalDOS8Dot3
- CloseEncryptedFileRaw
- CopyFile
- CopyFile2
- PCOPYFILE2\_PROGRESS\_ROUTINE
- CopyFileEx
- CopyFileTransacted

Download PDF

**CreateFile**

Creates or opens a file or I/O device. The most commonly used I/O devices are as follows: file, file stream, directory, physical disk, volume, console buffer, tape drive, communications resource, mailslot, and pipe.

**CreateFile2**

Creates or opens a file.

**CreateFileTransacted**

Creates or opens a file, file stream, or directory as a transacted operation.

**CreateHardLink**

Establishes a hard link between an existing file and a new file.

**CreateHardLinkTransacted**

Establishes a hard link between an existing file and a new file as a transacted operation.

**CreateIoCompletionPort**

Creates an input/output (I/O) completion port and associates it with a specified file handle, or creates an I/O completion port that is not yet associated with a file handle, allowing association at a later time.

**CreateSymbolicLink**

Creates a symbolic link.

Is this page helpful?

Yes No

In this article

[In this section](#)

# File Operations



File handling in C - javatpoint

javatpoint.com/file-handling-in-c

← → ↻ 🏠 🔒

☆ 📧 📅 27 🔄 🛠️ 🌐 ⋮

⊕ Type Casting

⊕ C Control Statement Test

✓ C Functions

⊕ What is function

⊕ Call: Value & Reference

⊕ Recursion in c

⊕ Storage Classes

⊕ C Functions Test

✓ C Array

⊕ 1-D Array

⊕ 2-D Array

⊕ Return an Array in C

⊕ Array to Function

⊕ C Array Test

✓ C Pointers

⊕ C Pointers

⊕ C Pointer to Pointer

⊕ C Pointer Arithmetic

⊕ Dangling Pointers in C

⊕ sizeof() operator in C

⊕ const Pointer in C

⊕ void pointer in C

⊕ C Dereference Pointer

⊕ Null Pointer in C

⊕ C Function Pointer

No.	Function	Description
1	fopen()	opens new or existing file
2	fprintf()	write data into the file
3	fscanf()	reads data from the file
4	fputc()	writes a character into the file
5	fgetc()	reads a character from file
6	fclose()	closes the file
7	fseek()	sets the file pointer to given position
8	fputw()	writes an integer to file
9	fgetw()	reads an integer from file
10	ftell()	returns current position
11	rewind()	sets the file pointer to the beginning of the file

GroupDocs

MultiBank Group


WORLD'S NO. 1 FX & CFD BROKER

PICK YOUR TRADING BONUS

START TRADING

ALL TRADING INVOLVES RISK.

# Operations Performed on Directory



- Search for a file
- Create a directory
- Delete a directory
- List a directory
- Rename a directory
- Traverse the file system



# Operations Performed on Directory



Directory Management Functions

docs.microsoft.com/en-us/windows/win32/fileio/directory-management-functions

Filter by title

functions

- Directory Management Functions
- CreateDirectory
- CreateDirectoryEx
- CreateDirectoryTransacted
- FindCloseChangeNotification
- FindFirstChangeNotification
- FindNextChangeNotification
- GetCurrentDirectory
- ReadDirectoryChangesExW
- ReadDirectoryChangesW

Download PDF

CreateDirectory	Creates a new directory.
CreateDirectoryEx	Creates a new directory with the attributes of a specified template directory.
CreateDirectoryTransacted	Creates a new directory as a transacted operation, with the attributes of a specified template directory.
FindCloseChangeNotification	Stops change notification handle monitoring.
FindFirstChangeNotification	Creates a change notification handle and sets up initial change notification filter conditions.
FindNextChangeNotification	Requests that the operating system signal a change notification handle the next time it detects an appropriate change.
GetCurrentDirectory	Retrieves the current directory for the current process.
ReadDirectoryChangesExW	Retrieves information that describes the changes within the specified directory, which can include extended information if that information type is specified.
ReadDirectoryChangesW	Retrieves information that describes the changes within the specified directory.

Is this page helpful?

Yes No

In this article

In this section

# Operations Performed on Directory



Reading a Directory | C For Dummies x +

c-for-dummies.com/blog/?p=3246

[Curl](#)

[Online Training](#)

[Bookstore](#)

I'm refusing to call it a "folder." That nonsense gained popularity with the Macintosh and then Windows. Before then, it was a *directory*, a list of files stored on media. Special C language functions are available to read and manipulate directories, which helps your programs manage files and do other fun file stuff.

A directory is really a special type of file, a data container that acts as a database referencing other files stored on the media. The media's file system determines how the files are organized and accessed. The directory holds all that information, such as the file's physical location, its name, timestamps, permissions, and other trivia. These details are accessible when you use the proper C language functions.

To access a directory, use the `opendir()` function. It's prototyped in the `dirent.h` header file as:

```
DIR *opendir(const char *filename);
```

The function requires a string argument, a name or path to a directory. The value returned is a *DIR* pointer, similar to the *FILE* pointer returned by `fopen()`.

After opening the directory and doing whatever, you use the `closedir()` function to close the directory:

```
int closedir(DIR *dirp);
```

The function requires a *DIR* pointer (`dirp`, `heh`) and returns 0 upon success, otherwise -1. Use the `errno` global variable to further examine the issue when -1 is returned.



# File/directory protection



# Protection

- File owner/creator should be able to control:
  - what can be done
  - by whom
- Types of access
  - Read
  - Write
  - Execute
  - Append
  - Delete
  - List


# Access Lists and Groups



- Mode of access: read, write, execute
- Three classes of users

			RWX
a) <b>owner access</b>	7	⇒	1 1 1
			RWX
b) <b>group access</b>	6	⇒	1 1 0
			RWX
c) <b>public access</b>	1	⇒	0 0 1

- Ask manager to create a group (unique name), say G, and add some users to the group.
- For a particular file (say *game*) or subdirectory, define an appropriate access.

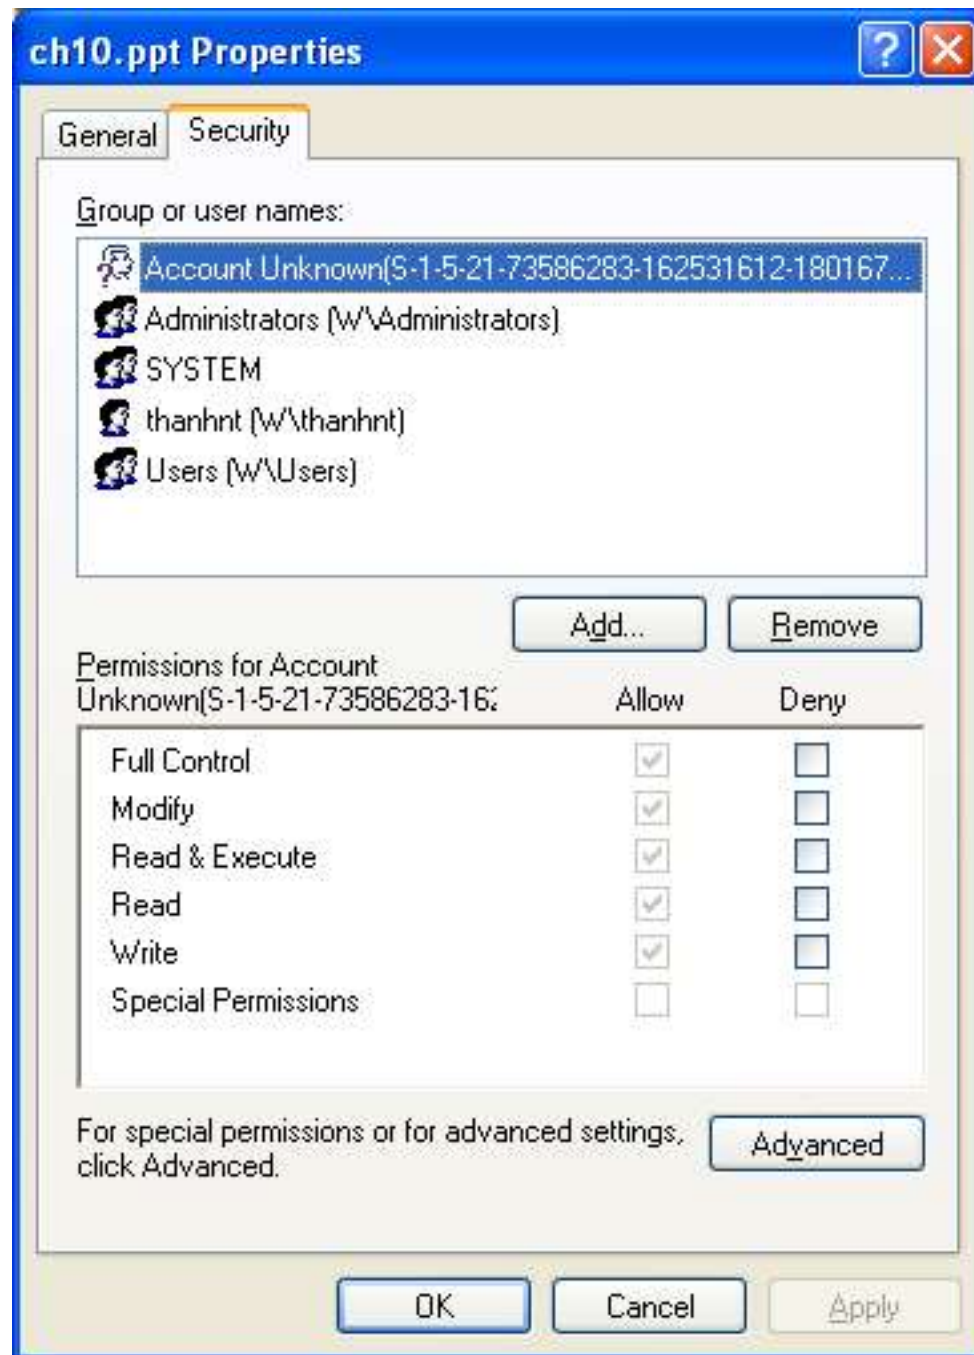


# Question



- Suppose a file has access mode 664. Which is correct ?
  - A. Any user can execute the file
  - B. Users of the owner group can execute the file
  - C. Any user can read the file
  - D. The owner cannot write the file

# Windows XP Access-control List Management





# Storage Allocation Methods



# Disk partitions



- A disk can be split into partitions
  - each is a consecutive range of cylinders
  - each is also called logic disk
  - e.g., Windows called: C, D, E drives

# Disk partitions



**PTDD Partition Table Doctor 3.5**

General Harddisk Partition Operations Sector View Help

Save Rebuild Restore Backup Fixboot Browse Check Undo Help

Harddisk 1 - 38166MB

FAT FAT32 NTFS Ext2/3 Swap Extend Free Other

1 (C:) winXP HPFS/NTFS 9546MB

2 (E:) play HPFS/NTFS 4996MB

3 (F:) music HPFS/NTFS 5004MB

4 (G:) SOFT FAT16 454MB

5 (H:) PLUG FAT16 250MB

[Free] 9287MB

[Free] 8620MB


Partition	Active	System	Label	Start			Space(MB)
				C	H	S	
1 <C:>	Yes	HPFS/NTFS	winXP	0	1	1	9546
	No	EXTEND		1217	0	1	19994
	No	[Free]		3766	0	1	8620
<b>Logic Partition</b>							
2 <E:>	No	HPFS/NTFS	play	1217	1	1	4996
3 <F:>	No	HPFS/NTFS	music	1854	1	1	5004
4 <G:>	No	FAT16	SOFT	2492	1	1	454
5 <H:>	No	FAT16	PLUG	2550	1	1	250
	No	[Free]		2582	1	1	9287

We recommend closing all other applications while running Partition Table Doctor.

ns  
ylinders

ves

# Disk partitions



PTDD Partition Table Doctor 3.5

General Harddisk Partition Operations Sector View Help

Save Rebuild Restore Backup Fixboot Browse Check Undo Help

Harddisk 1 - 38166MB

1 (C:) winXP HPFS/NTFS 9546MB

2 (E:) play HPFS/NTFS 4936MB

3 (F:) music HPFS/NTFS 5004MB

4 (G:) SOFT FAT16 2492MB

5 (H:) PLUG FAT16 2550MB

[Free] 9287MB


[Free] 8620MB

Partition	Active	System	Label	Start	C	H	S
1 <C:>	Yes	HPFS/NTFS	winXP		0	1	1
	No	EXTEND		1217	0	1	
	No	[Free]		3766	0	1	

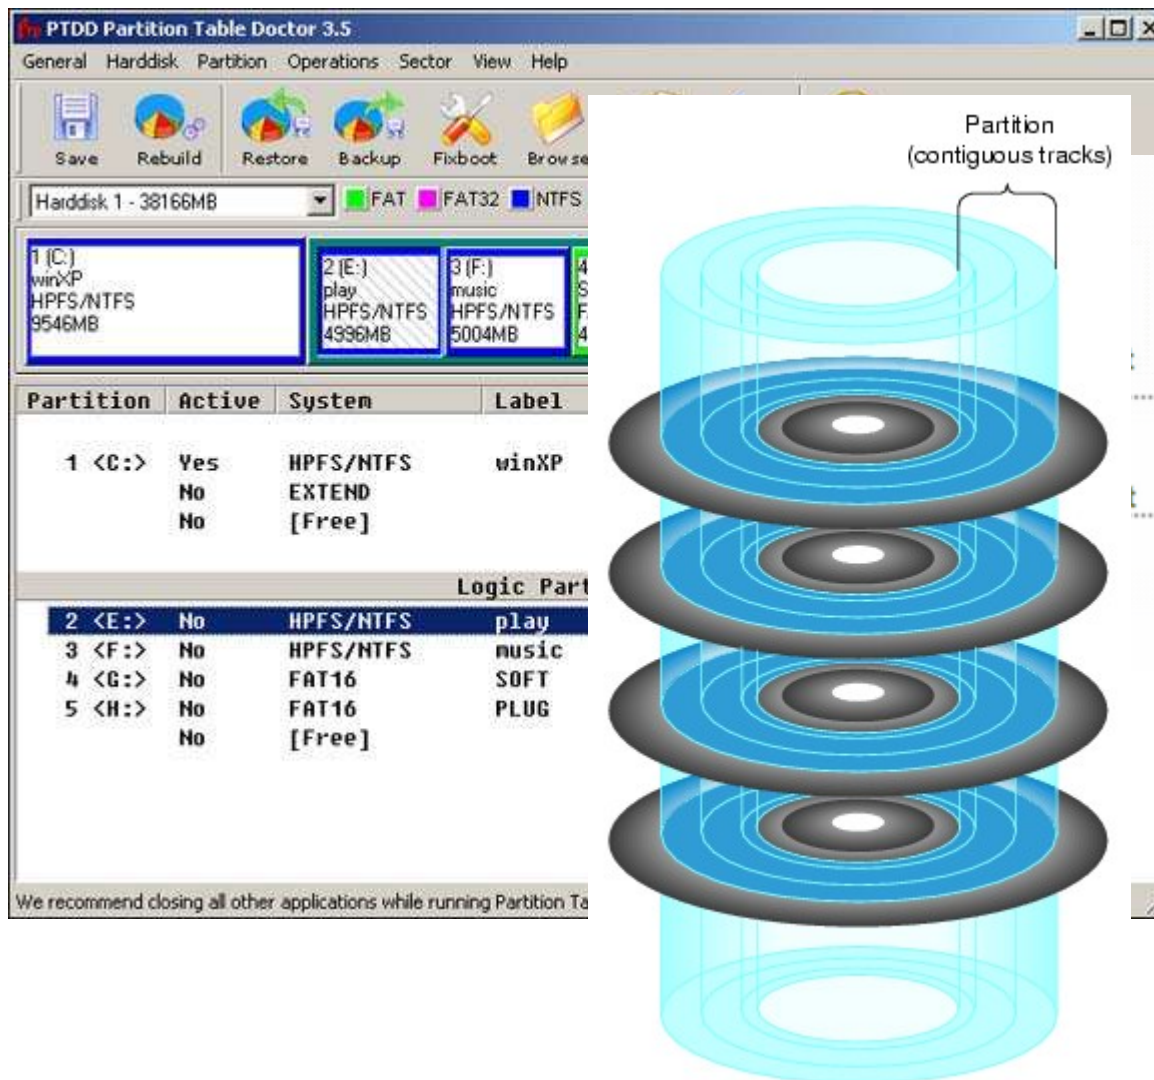
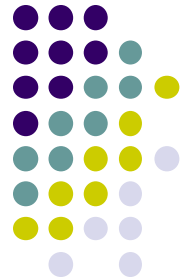
Logic Partition

Partition	Active	System	Label	Start	C	H	S
2 <E:>	No	HPFS/NTFS	play	1217	1	1	
3 <F:>	No	HPFS/NTFS	music	1854	1	1	
4 <G:>	No	FAT16	SOFT	2492	1	1	454
5 <H:>	No	FAT16	PLUG	2550	1	1	250
	No	[Free]		2582	1	1	9287


We recommend closing all other applications while running Partition Table Doctor.



# Disk partitions




ns



# Partition organization

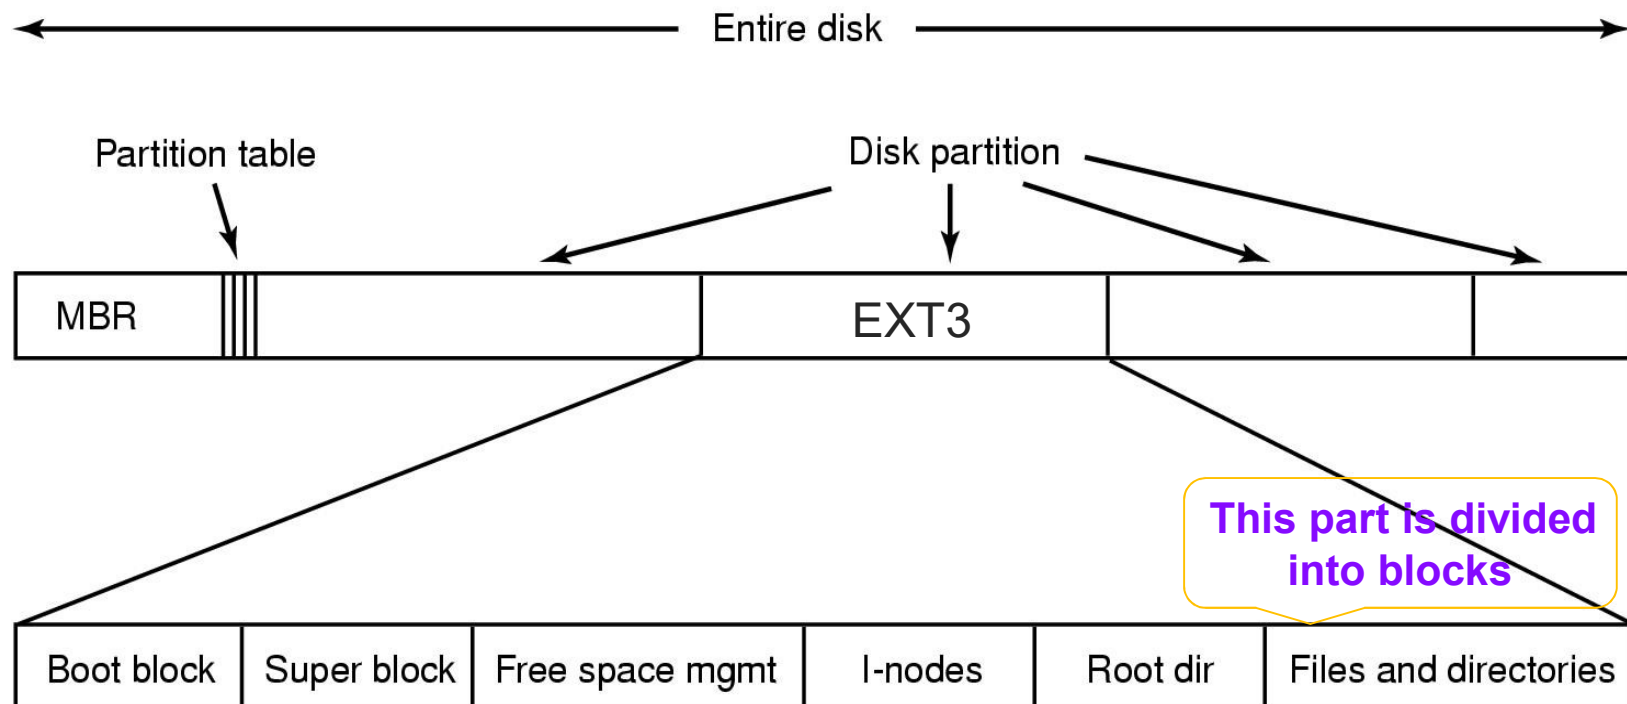


- Organization is specific to each OS
  - region for storing data is divided into equal **blocks**
  - block is a read/write unit of OS




# Partition organization

- Organization is specific to each OS
  - region for storing data is divided into equal **blocks**



# Format logical volume



The image shows a Windows File Explorer window with a context menu open for the drive 'Mars (D:)'. The drive is shown as a USB icon with a capacity of 931 GB and 801 GB free. The context menu includes options like 'Open', 'Scan with Microsoft Defender...', 'Upload to MEGA', 'Give access to', 'Restore previous versions', '7-Zip', 'CRC SHA', 'Include in library', 'Pin to Start', 'Format...', 'Copy', 'Create shortcut', 'Rename', and 'Properties'. The 'Format...' option is highlighted with a red rectangle.

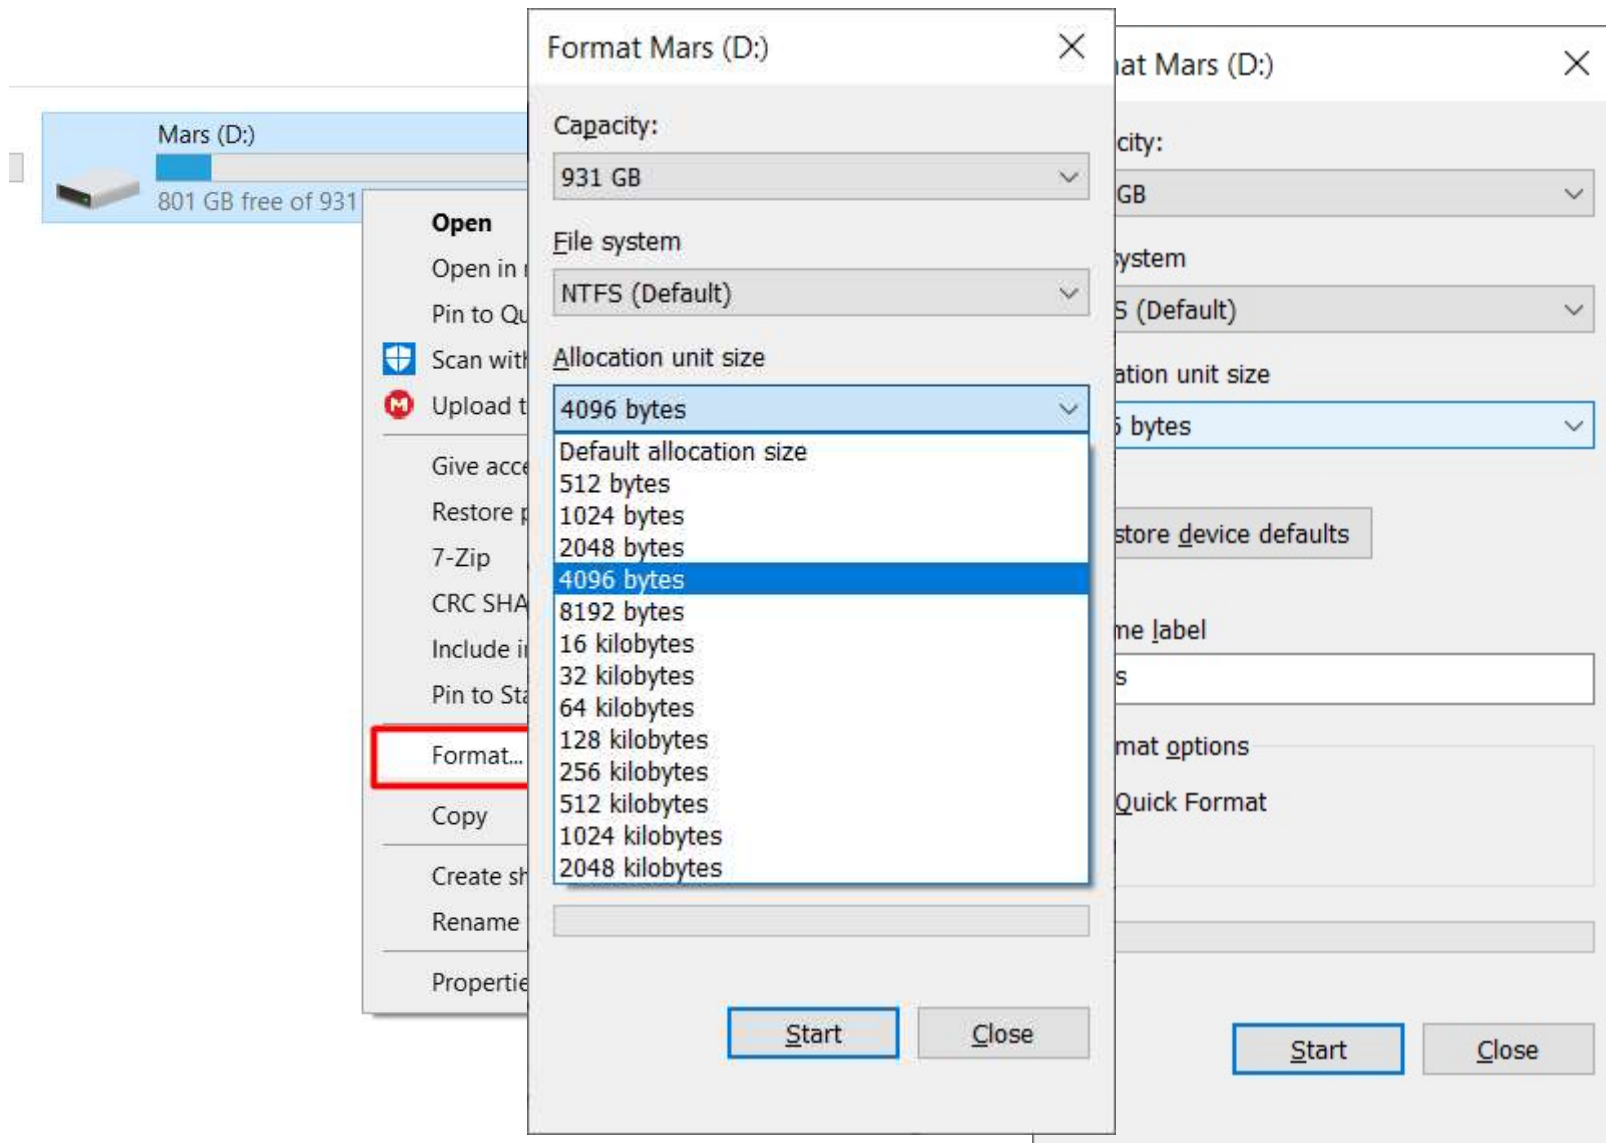
To the right of the File Explorer is the 'Format Mars (D:)' dialog box. It contains the following settings:

- Capacity: 931 GB
- File system: NTFS (Default)
- Allocation unit size: 4096 bytes
- Restore device defaults button
- Volume label: Mars
- Format options: ☒ Quick Format
- Start and Close buttons at the bottom.





# Format logical volume





# xfs format



mkfs.xfs(8)

System Manager's Manual


mkfs.xfs(8)

## NAME [top](#)

mkfs.xfs - construct an XFS filesystem

## SYNOPSIS [top](#)

```
mkfs.xfs [ -b block size options ] [ -m global_metadata_options ]  
[ -d data_section_options ] [ -f ] [ -i inode_options ] [ -l  
log_section_options ] [ -n naming_options ] [ -p protofile ] [ -q  
] [ -r realtime_section_options ] [ -s sector_size_options ] [ -L  
label ] [ -N ] [ -K ] device  
mkfs.xfs -V
```



# ext4 format

## mkfs.ext4(8) - Linux man page

### Name

mkfs.ext4 - create an ext2/ext3/ext4 filesystem

### Synopsis

```
mkfs [ -c | -l filename ] [ -b block-size ] [ -f  
fragment-size ] [ -g blocks-per-group ] [ -G  
number-of-groups ] [ -i bytes-per-inode ] [ -I  
inode-size ] [ -j ] [ -J journal-options ] [ -K ] [ -  
N number-of-inodes ] [ -n ] [ -m reserved-  
blocks-percentage ] [ -o creator-os ] [ -O  
feature[,...] ] [ -q ] [ -r fs-revision-level ] [ -E  
extended-options ] [ -v ] [ -F ] [ -L volume-  
label ] [ -M last-mounted-directory ] [ -S ] [ -t  
fs-type ] [ -T usage-type ] [ -U UUID ] [ -V ]  
device [ blocks-count ]
```

```
mkfs -O journal_dev [ -b block-size ] [ -L  
volume-label ] [ -n ] [ -q ] [ -v ] external-  
journal [ blocks-count ]
```

# Allocation Methods



- An allocation method refers to how disk blocks are allocated for files
  - Contiguous allocation
  - Linked allocation
  - Indexed allocation
- Each method needs an appropriate way to access file

# Contiguous Allocation



- Each file occupies a set of contiguous blocks on the disk
- Simple
  - only starting location (block #) and length (number of blocks) are required
- Random access
- Wasteful of space
  - dynamic storage-allocation problem
- Files cannot grow

# Mapping from logical add into block/offset



0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	...
0	1	2	3	0	1	2	3	0	1	2	3	0	1	2	3	...
0				1				2				3				

- Logical address = x
- $Q \text{ (blockid)} = x / \text{block\_size}$
- $R \text{ (offset)} = x \% \text{block\_size}$
- $X = 9 \Rightarrow Q = 9/4 = 2; R = 9\%4 = 1$

# Contiguous Allocation



- Mapping from logical to physical

- LA is position to access


- - LA/512 — Q - block
    - R - offset

Suppose block size is 512KB

Block to be accessed =  $Q + \text{starting address}$

Displacement/offset into block =  $R$

# Contiguous Allocation of Disk Space





## Question

- A system using contiguous allocation
  - block size is 2KB
  - a file has the length of 12.5MB
- Which of the following is the correct location of file at position 50.5KB
  - A. block 25, offset 512
  - B. block 24, offset 1024
  - C. block 25, offset 510
  - D. block 24, offset 2512





# Extent-Based Systems

- Many newer file systems (I.e. Veritas File System)
  - use a modified contiguous allocation scheme
- Extent-based file systems
  - allocate disk blocks in **extents**
- An **extent** is a contiguous blocks of disk
  - Extents are allocated for file allocation
  - A file consists of one or more extents
  - extents of a file are not necessarily contiguous



## Question

- An extent-based file system
  - an extent has *100* blocks
  - a block has size *2KB*
  - a file has size *25.3MB*
- Which of the following is the correct extent number (started by 0) of file at position 15MB?
  - A. 74
  - B. 75
  - C. 76
  - D. 77



## Question

- An extent-based file system
  - an extent has *100* blocks
  - a block has size *2KB*
  - a file has size *25.3MB*
- Which of the following is the correct block (started by 0) number in the extent containing the data of file at position *15MB*?
  - A. 77
  - B. 78
  - C. 79
  - D. 80



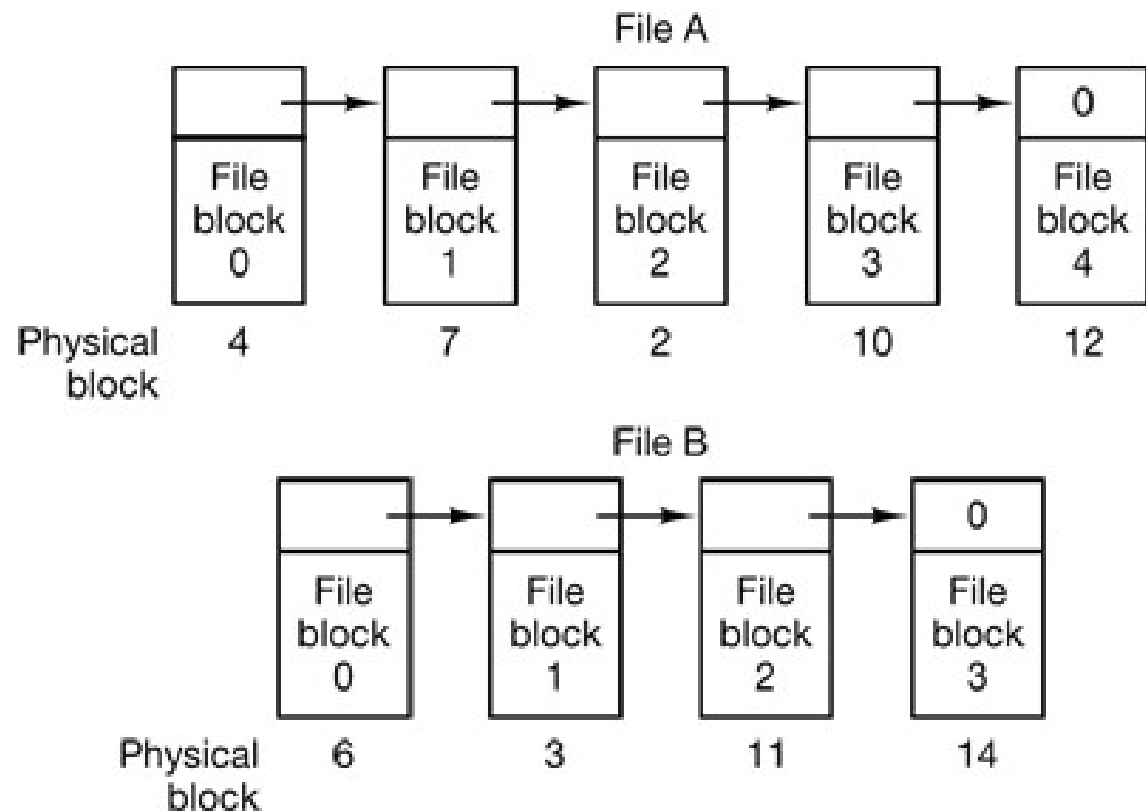

## Question

- An extent-based file system
  - an extent has *100* blocks
  - a block has size *2KB*
  - a file has size *25.3MB*
- Which of the following is the correct offset of the of file at position *15MB* in the block containing data?
  - A. 0
  - B. 2047
  - C. 2048
  - D. 512


# Linked Allocation



- Each file is a linked list of disk blocks:
  - blocks may be scattered anywhere on the disk.



# Linked Allocation

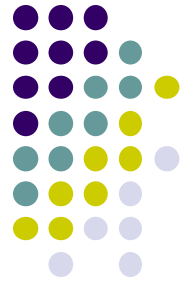


# Question



- A system uses linked list allocation
  - partition size 500GB (1GB=1024MB, 1MB=1024KB,...)
  - block size 1KB
- Which of the following is the correct size of the pointer of each block?
  - A. short (2 bytes)
  - B. int (4 bytes)
  - C. float (4 bytes)
  - D. long (8 bytes)

# Question



- A system uses linked list allocation
- Which of the following is the correct reason for **no direct access**?
  - A. because we don't know the location of block  $n$  directly
  - B. because data blocks of a file may be scattered
  - C. because of security reason
  - D. because the information of data block location is hidden



# Question



- A system uses linked list allocation
- Which of the following is the correct way to access block  $n$  of a file?
  - A. read the first block to find the location of block  $n$
  - B. look up the location of block  $n$  from a table
  - C. recursively read block  $n-1$  to find out the location of block  $n$
  - D. there is no way to find the location of block  $n$

# Linked Allocation (Cont.)



- Simple
  - need only starting address
- Free-space management system
  - no waste of space
- No random access
  - File-allocation table (FAT) – disk-space allocation used by MS-DOS and OS/2.

# Question



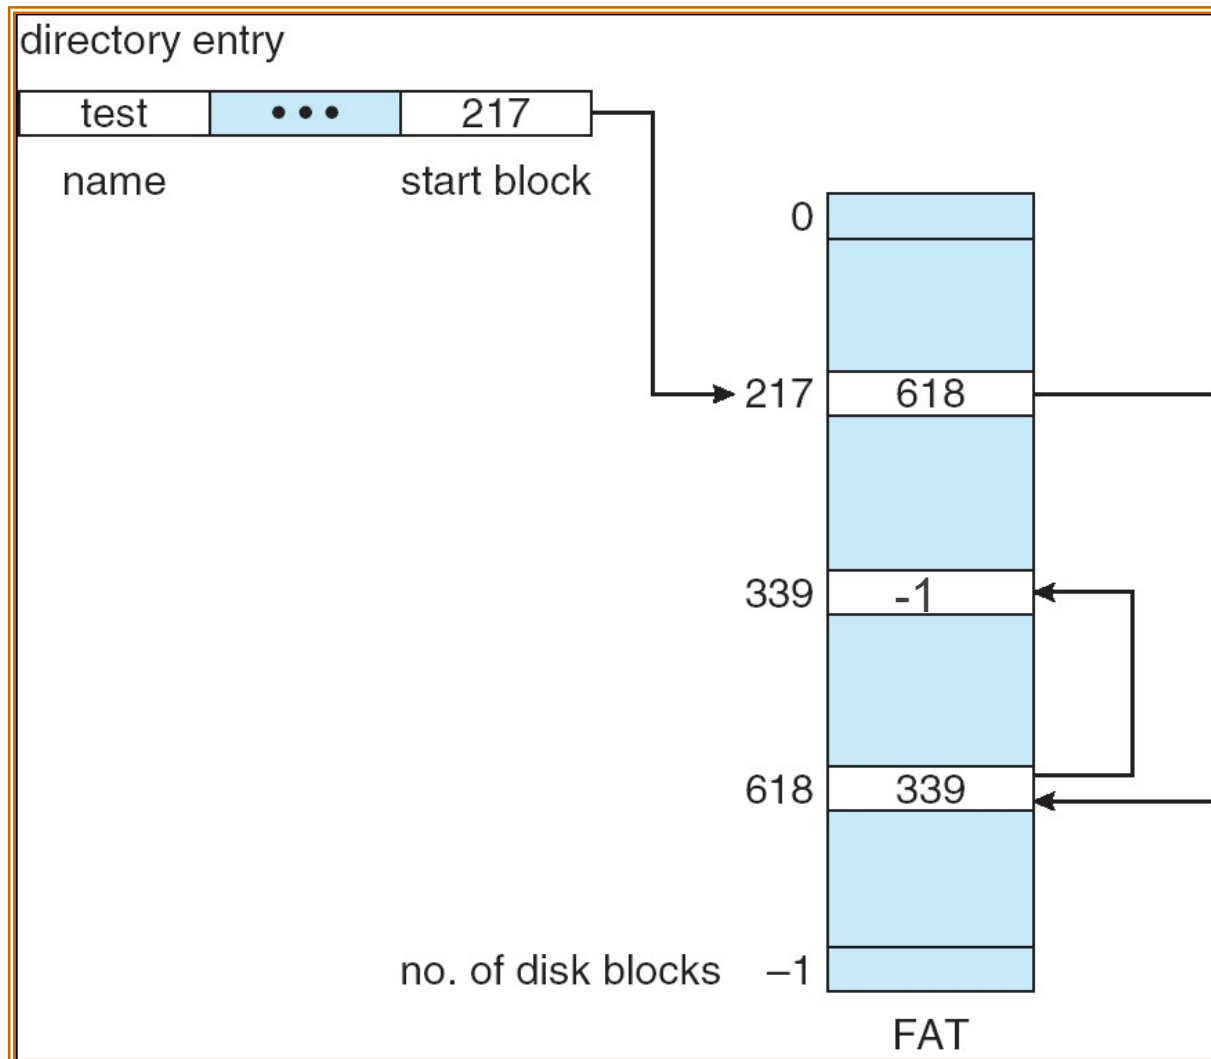
- A system uses linked list allocation
  - block size 2KB
  - pointer size 4 bytes
  - file size 15.4MB
- Which of the following is the correct block of file at position 15.25KB?
  - A. 7
  - B. 8
  - C. 9
  - D. 10

# Question



- A system uses linked list allocation
  - block size 2KB
  - pointer size 4 bytes
  - file size 15.4MB
- Which of the following is the correct offset in the block of file at position 15.25KB?
  - A. 1311
  - B. 1312
  - C. 1313
  - D. 1314

# File-Allocation Table (DOS)



- Separate the pointers from data
- pointers are stored in File Allocation Table (FAT)
- The system stores two identical copies of FAT

# Question




- Which of the following is incorrect about FAT?
  - A. it is fast to access the block  $n$  of a file
  - B. it is slow to access the block  $n$  of a file
  - C. if FAT is corrupted the whole partition is corrupted
  - D. the system keeps two copies of FAT in order to reduce the risk of FAT corruption

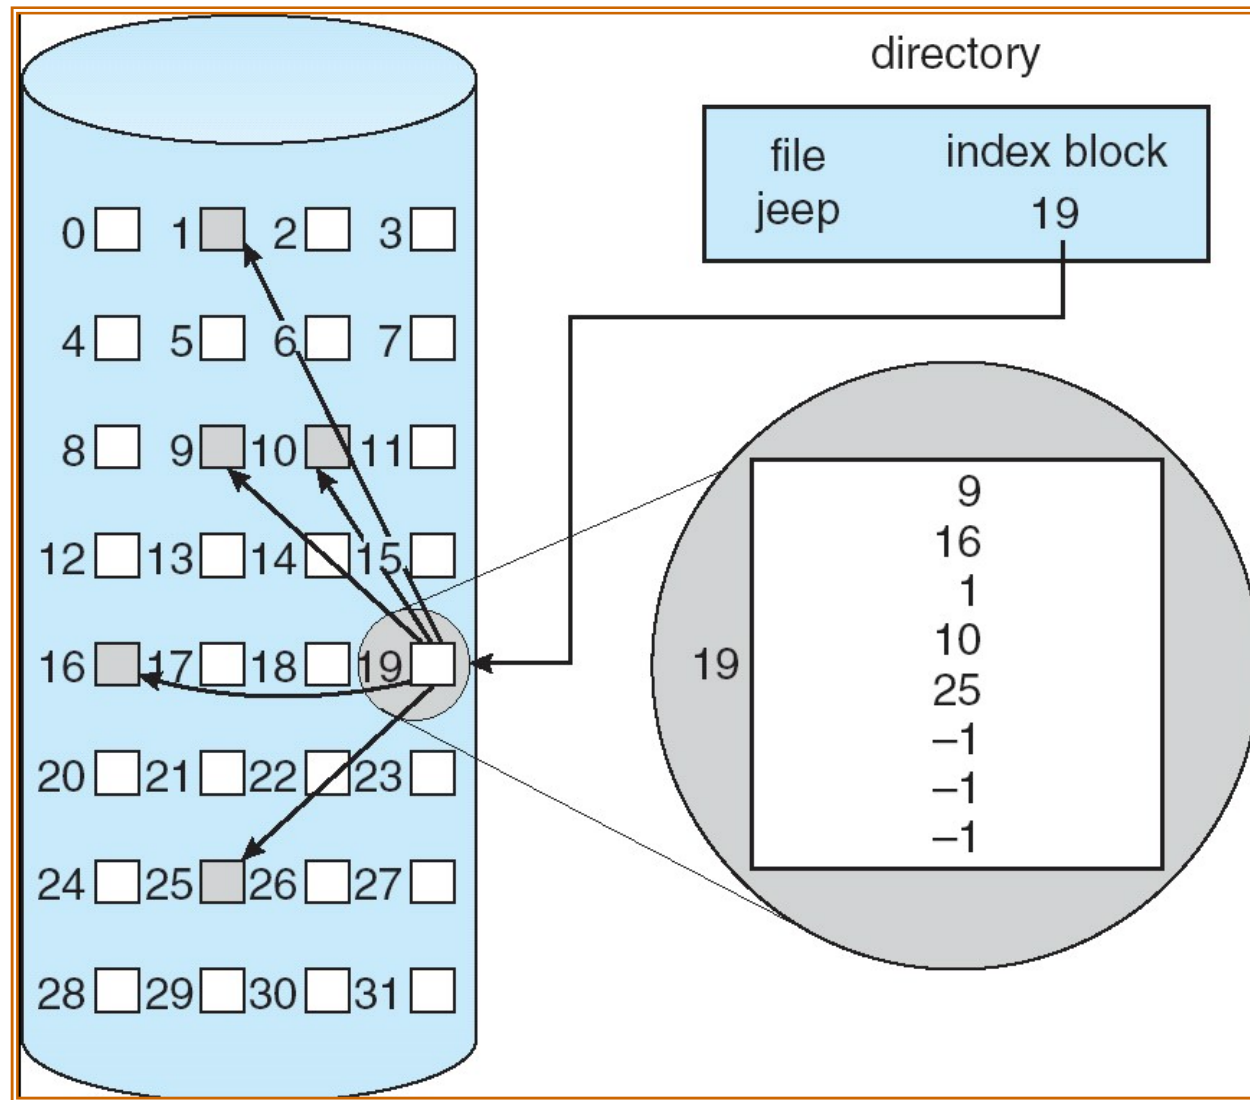

# Indexed Allocation



- Brings all pointers together into the *index block*
- Logical view



# Example of Indexed Allocation





# Indexed Allocation (Cont.)

- Need index block
- Random access
- Dynamic allocation without external fragmentation
  - have overhead of index block



# Question



- System uses indexed allocation
  - block size 4KB
  - pointer size 4 bytes
- Which of the following is the correctly maximum file size?
  - A. 4MB
  - B. 8MB
  - C. 16MB
  - D. 32MB




## Question

- System uses indexed allocation
  - block size 4KB
  - pointer size 4 bytes
  - file size 3MB
- Which of the following is the correct block (starting from 0) and offset at file position 35KB?
  - A.  $(block, offset) = (9, 3071)$
  - B.  $(block, offset) = (9, 3070)$
  - C.  $(block, offset) = (8, 3072)$
  - D.  $(block, offset) = (8, 3070)$

# Indexed Allocation (cont'd)



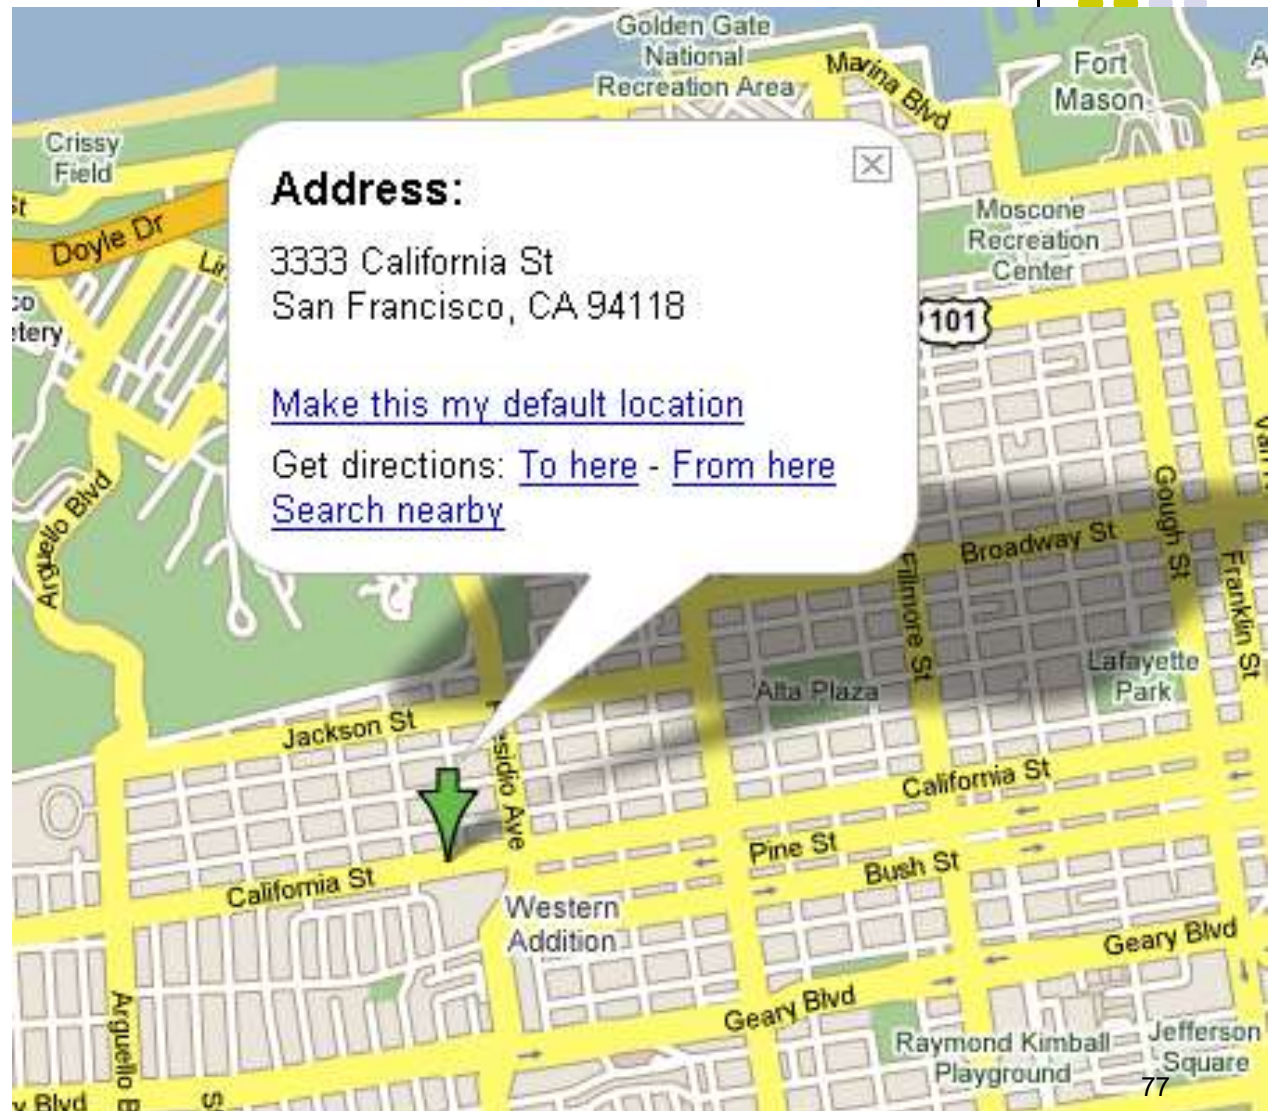
- Two-level index
  - Two level of index block
  - Method to increase file size



# Address locating



Go to **USA** →  
**San Francisco**  
→  
**California St**  
→ **3333**



# Indexed Allocation (cont'd)



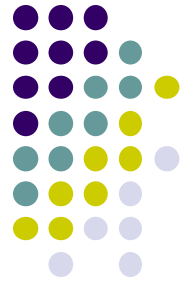
- A system uses two-level index
  - block size 4KB
  - pointer size 4 Bytes
- Which of the following is the correct maximum file size?
  - A. 4GB
  - B. 1GB
  - C. 8GB
  - D. 2GB

# Indexed Allocation (cont'd)



- A system uses two-level index
- Which of the following is the correct steps to locate the data at file position  $n$ ?
  - A. Identify block number → block number in block table → offset
  - B. Identify block number in outer index block → block number → offset
  - C. Identify offset → block number
  - D. none of the above

# Indexed Allocation (cont'd)




- A system uses two-level index
  - block size 4KB
  - pointer size 4 Bytes
  - File size 20MB
- Which is the correct address of file at position 15MB?
  - A. (4,3072,0)
  - B. (3,768,1023)
  - C. (3,768,0)
  - D. (3,3072,0)



# Indexed Allocation (Cont.)

- Linked scheme
  - no limit on file size
- Combine linked list with index block
  - index blocks are linked
  - last pointer of a index block is the address of the next one



index table

data blocks



# Question



- Linked index block allocation
  - block size 2KB
  - pointer size 4 bytes
  - File size 20MB
- Which of the following is the correct **steps** to read file at position 15.5MB
  - A. identify index block → offset → block number → read
  - B. identify offset → block number → index block → read
  - C. identify offset → index block → block number → read
  - D. identify index block → block number → offset → read

# Question



- Linked index block allocation
  - block size 2KB
  - pointer size 4 bytes
  - File size 20MB
- Which of the following is the correct **index block** number at file position 15.5MB (start from 0)

# Question

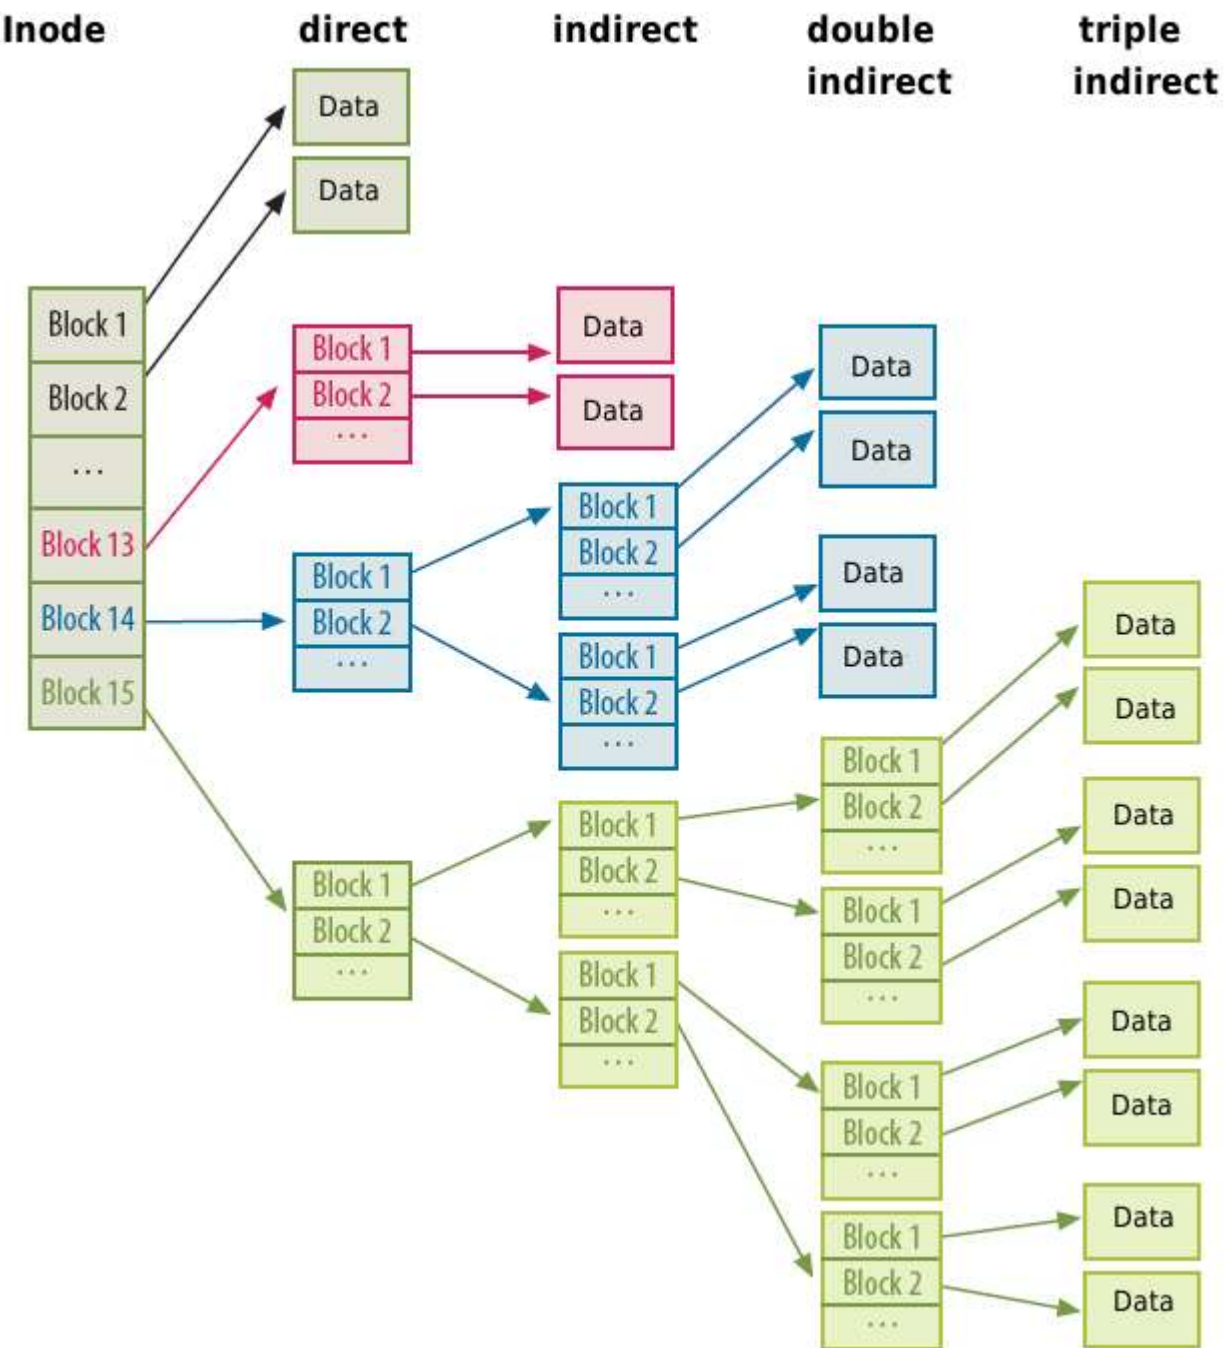



- Linked index block allocation
  - block size 2KB
  - pointer size 4 bytes
  - File size 20MB
- Which of the following is the correct **index block** number at file position 15.5MB (start from 0)
  - A. 13
  - B. 14
  - C. 15
  - D. 16

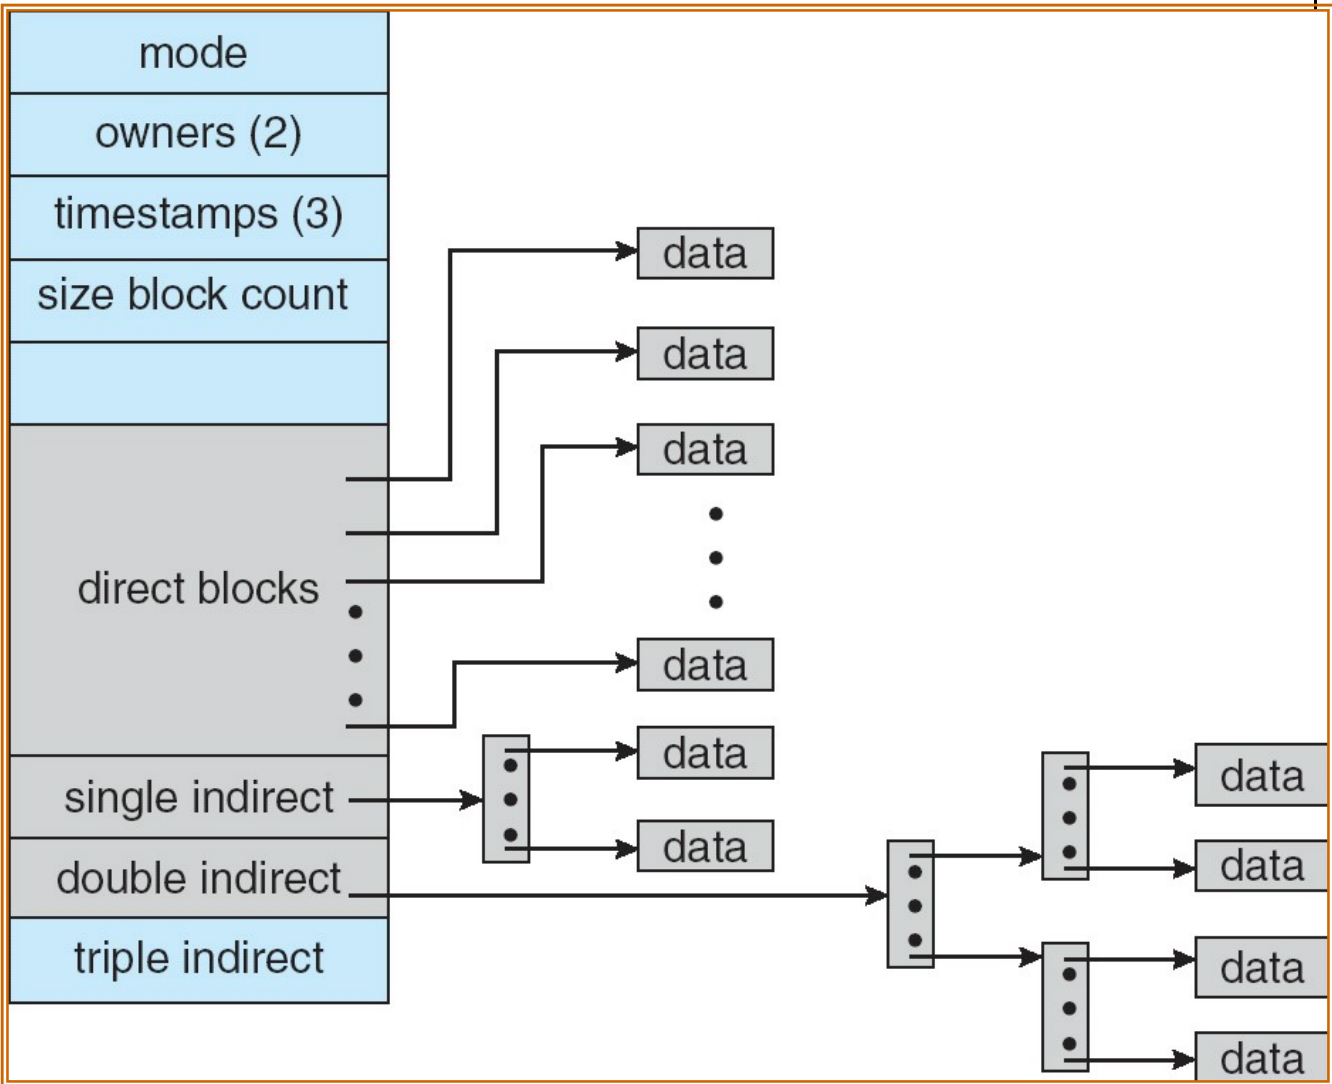


## Question

- Linked index block allocation
  - block size 2KB
  - pointer size 4 bytes
  - File size 20MB
- Which of the following is the correct **block** number and **offset** at file position 15.5MB
  - A.  $(block, offset) = (271, 2047)$
  - B.  $(block, offset) = (271, 0)$
  - C.  $(block, offset) = (270, 2047)$
  - D.  $(block, offset) = (270, 0)$



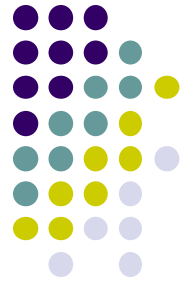
**Combined Scheme: UNIX**  
**(4K bytes per block)**



# LINUX allocation

(10 direct pointers)

# Question



- A UNIX system
  - pointer size 4 bytes
  - block size 4 KB
  - 12 direct pointers, 1 single indirect, 1 double indirect, 1 triple indirect pointers
- Which of the following is the correct **maximum file size**?
  - A.  $(12+2^{10}+2^{20}+2^{30})\text{KB}$
  - B.  $4*(2^{10}+2^{20}+2^{30})\text{KB}$
  - C.  $2^{32}\text{KB}$
  - D.  $4*(12+2^{10}+2^{20}+2^{30})\text{KB}$



# Question



- A UNIX system
  - pointer size 4 bytes
  - block size 4 KB
  - 12 direct pointer, 1 single indirect, 1 double indirect, 1 triple indirect pointers
- Which of the following is the correct **maximum number of indexed blocks**?
  - A.  $(1+2^{10}+2^{20})$
  - B.  $(2+2^{10}+2^{20})$
  - C.  $(3+2^{11}+2^{20})$
  - D.  $2^{20}$

# Question




- A UNIX system
  - pointer size 4 bytes
  - block size 4 KB
  - 12 direct pointer, 1 single indirect, 1 double indirect, 1 triple indirect pointers
  - File size 78MB
- Which of the following is the correct **location** at file position 95KB?
  - A. triple indirect block,  $(block, offset)=(12,3071)$
  - B. double indirect block,  $(block, offset)=(12,3071)$
  - C. single indirect block,  $(block, offset)=(11,3071)$
  - D. single indirect block,  $(block, offset)=(11,3072)$

# Free-Space Management



- Bit vector ( $n$  blocks), e.g. Linux (ext3)
  - Easy to get contiguous blocks



$$\text{bit}[i] = \begin{cases} 0 \Rightarrow \text{block}[i] \text{ free} \\ 1 \Rightarrow \text{block}[i] \text{ occupied} \end{cases}$$

- Bit map requires extra space

Example:

block size =  $2^{12}$  bytes

disk size =  $2^{30}$  bytes (1 gigabyte)

$n = 2^{30}/2^{12} = 2^{18}$  bits (or 32K bytes)

# Free-Space Management (Cont.)




- Linked list (free list)
  - Cannot get contiguous space easily
  - No waste of space
- Grouping
  - Use linked blocks to store pointers to free blocks
  - Last pointer in a block is the address of the next one
- Counting
  - several contiguous blocks are freed/allocated for a file
  - each entry has
    - address of the first free block
    - number of contiguously free blocks

# Free-Space Management (Cont.)



- Need to protect:
  - Pointer to free list
  - Bit map
    - Must be kept on disk
    - Copy in memory and disk may differ
    - Cannot allow for block[ $i$ ] to have a situation where  $\text{bit}[i] = 1$  in memory and  $\text{bit}[i] = 0$  on disk
  - Solution:
    - Set  $\text{bit}[i] = 1$  in disk
    - Allocate block[ $i$ ]
    - Set  $\text{bit}[i] = 1$  in memory

# Linked Free Space List on Disk







# Shared file systems

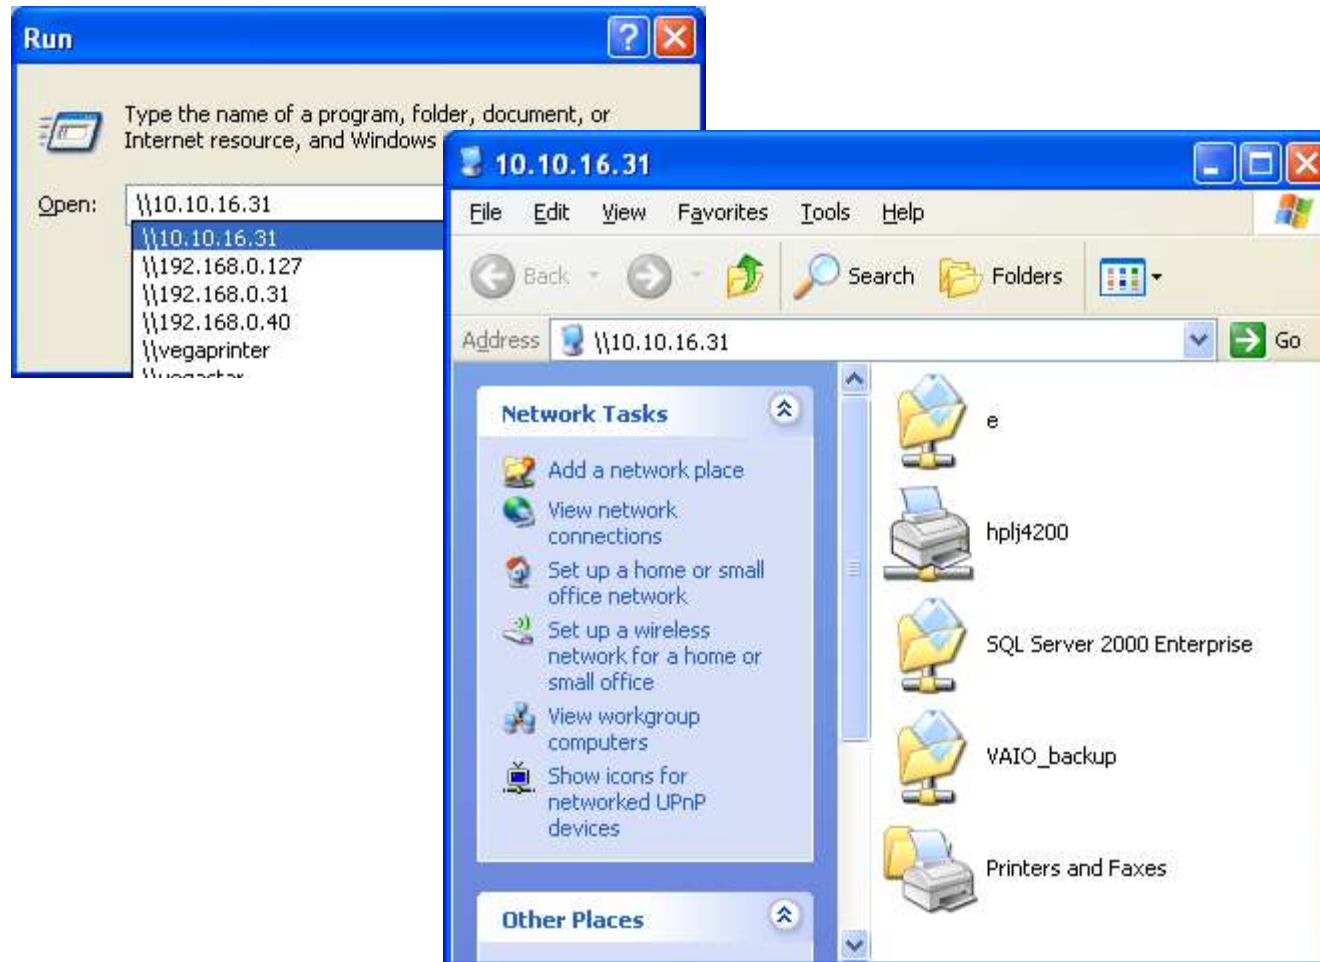
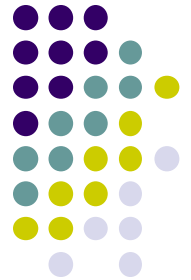
- A file system is shared to other machines
  - the file system may be large and powerful
  - file sharing is needed in many applications
  - available in many systems, e.g., Windows, Linux (NFS)

# Shared file systems

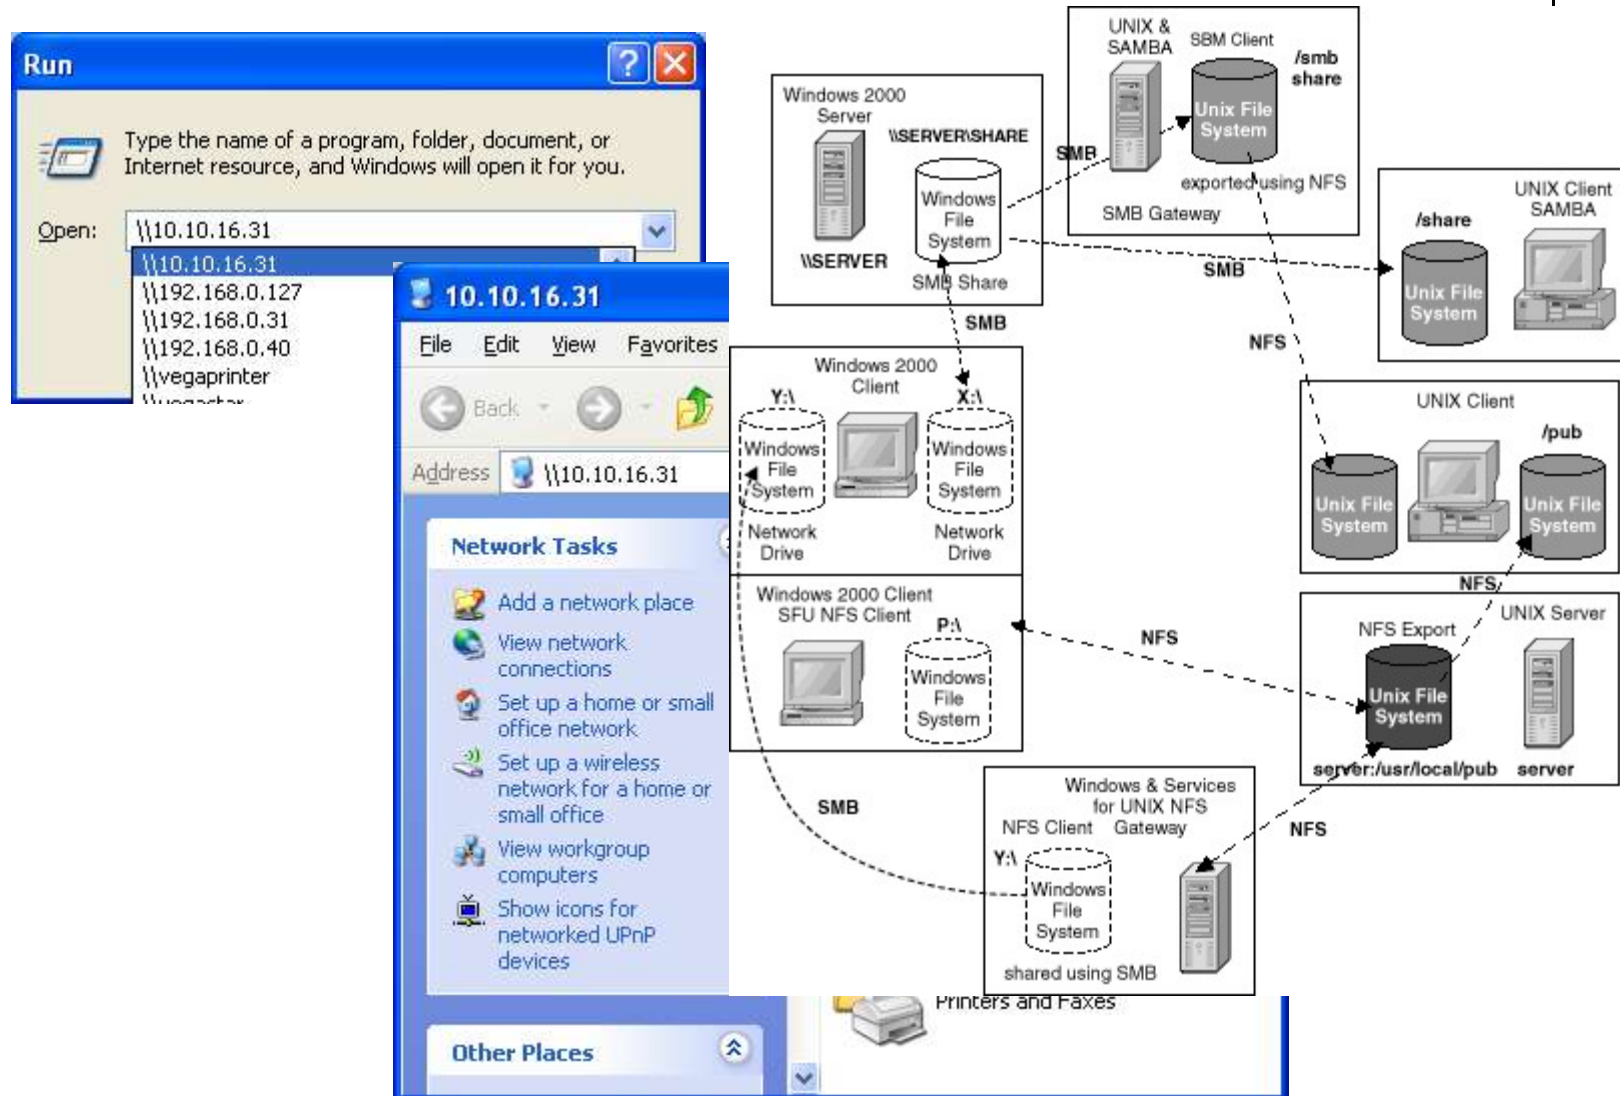




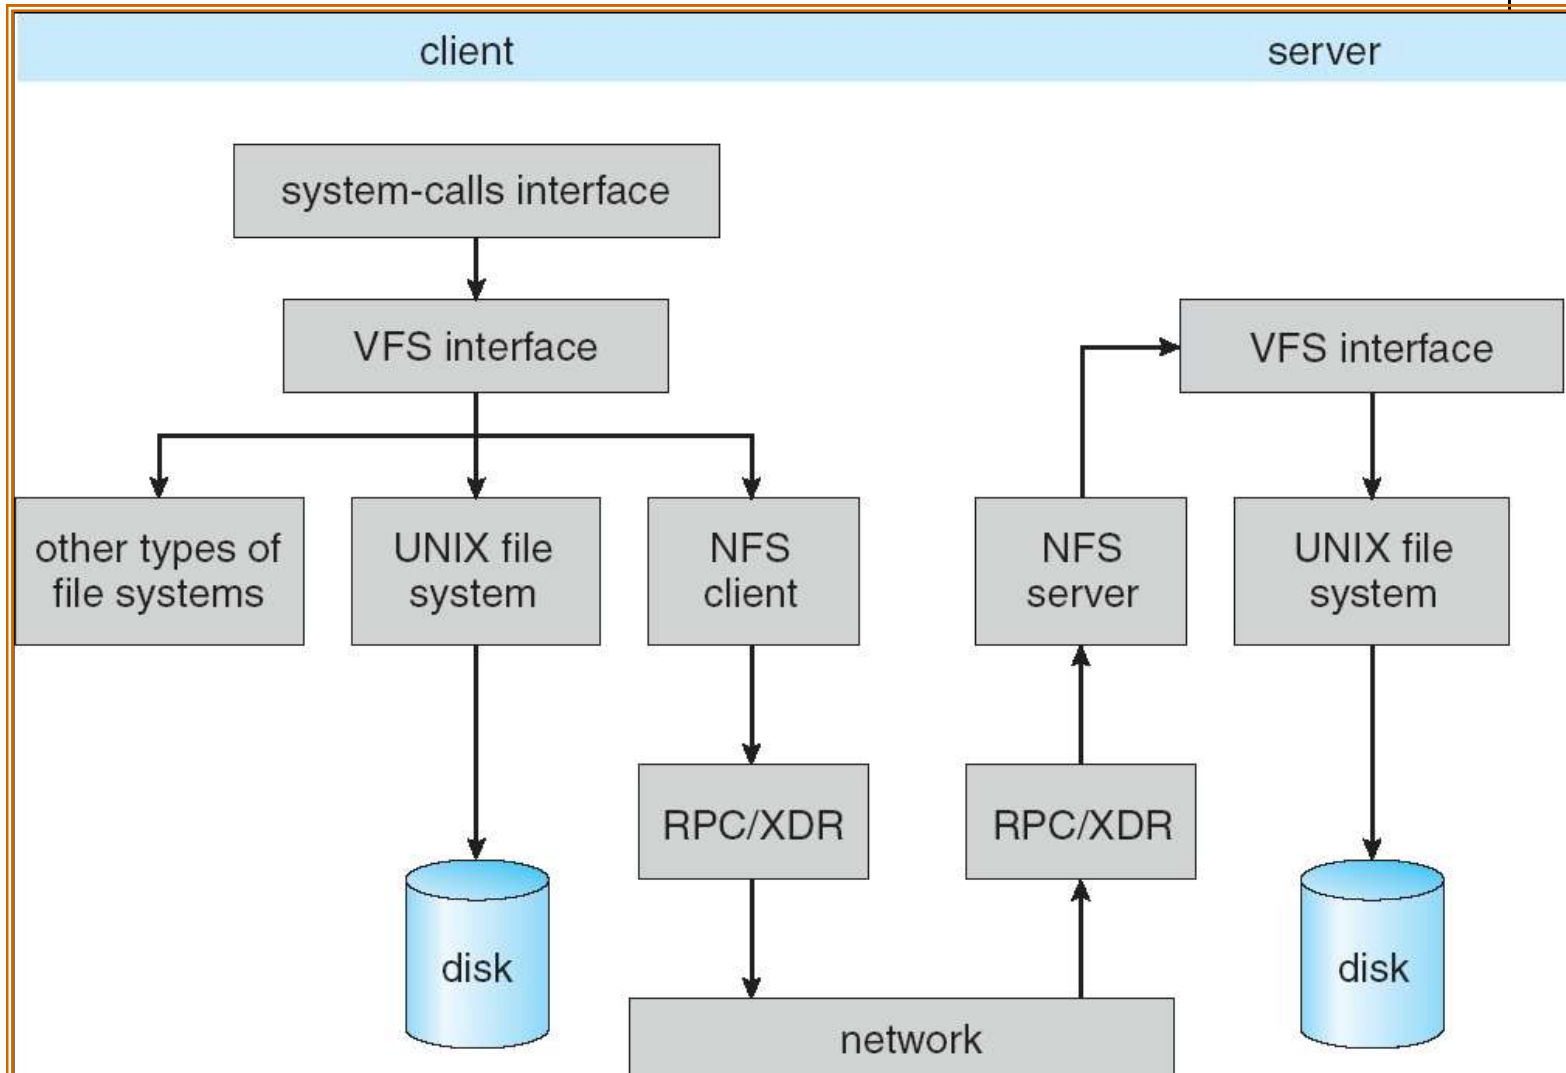
# Shared file systems



# Shared file systems



# Network File System (NFS) Architecture





**Question?**